

# Computer Architecture

## The P6 Microarchitecture

An Example of an Out-Of-Order Micro-processor

By Dan Tsafrir, 2/5/2011, 16/5/2011

Presentation based on slides by Lihu Rappoport and Adi Yoaz

# The P6 Family (i686)

- **Features**

- 1<sup>st</sup> out of order x86 (=> data flow analysis)
- Speculative execution (across branches; requires flush+recovery)
- Multiple branch prediction (wide op window contains 5 branch on avg)
- Register renaming (solves false dependencies, gives more regs)
- Super-pipeline: ~12 pipe stages (P-IV had 31! i7 back to 14)

Processor	Year	Freq (MHz)	Bus (MHz)	L2 cache	Feature size**
Pentium® Pro	1995	150~200	60/66	256/512K*	0.5, 0.35µm
Pentium® II	1997	233~450	66/100	512K*	0.35, 0.25µm
Pentium® III	1999	450~1400	100/133	256/512K	0.25, 0.18, 0.13µm
Pentium® M	2003	900~2260	400/533	1M / 2M	0.13, 90nm
Core™	2005	1660~2330	533/667	2M	65nm
Core™ 2	2006	1800~2930	800/1066	2/4/8M	65nm

\*off die

\*\* size of smallest part is smaller than the feature size

# The P6 Family (i686)



- Still used:
  - MacBook Air (1.4GHz Core 2 Duo)
  - Good for low power consumption
- Clock frequency ~proportional to feature size
- After P-III came P-IV... which wasn't ideal for mobile computing
- Much (not all) of the improvement comes from feature size minimization

Processor	Year	Freq (MHz)	Bus (MHz)	L2 cache	Feature size**
Pentium® Pro	1995	150~200	60/66	256/512K*	0.5, 0.35µm
Pentium® II	1997	233~450	66/100	512K*	0.35, 0.25µm
Pentium® III	1999	450~1400	100/133	256/512K	0.25, 0.18, 0.13µm
Pentium® M	2003	900~2260	400/533	1M / 2M	0.13, 90nm
Core™	2005	1660~2330	533/667	2M	65nm
Core™ 2	2006	1800~2930	800/1066	2/4/8M	65nm

\*off die

\*\* size of smallest part is smaller than the feature size

# Chip logically partitioned to 3

- **Front end**

- In order, get and ops from memory
- Decode them + turn them from CISC ops to  $\geq 1$  u-ops
- Uops are RISC-like
- So x86 input=CISC, but internally it's RISC
- The front-end is responsible to make the transition

- **Core**

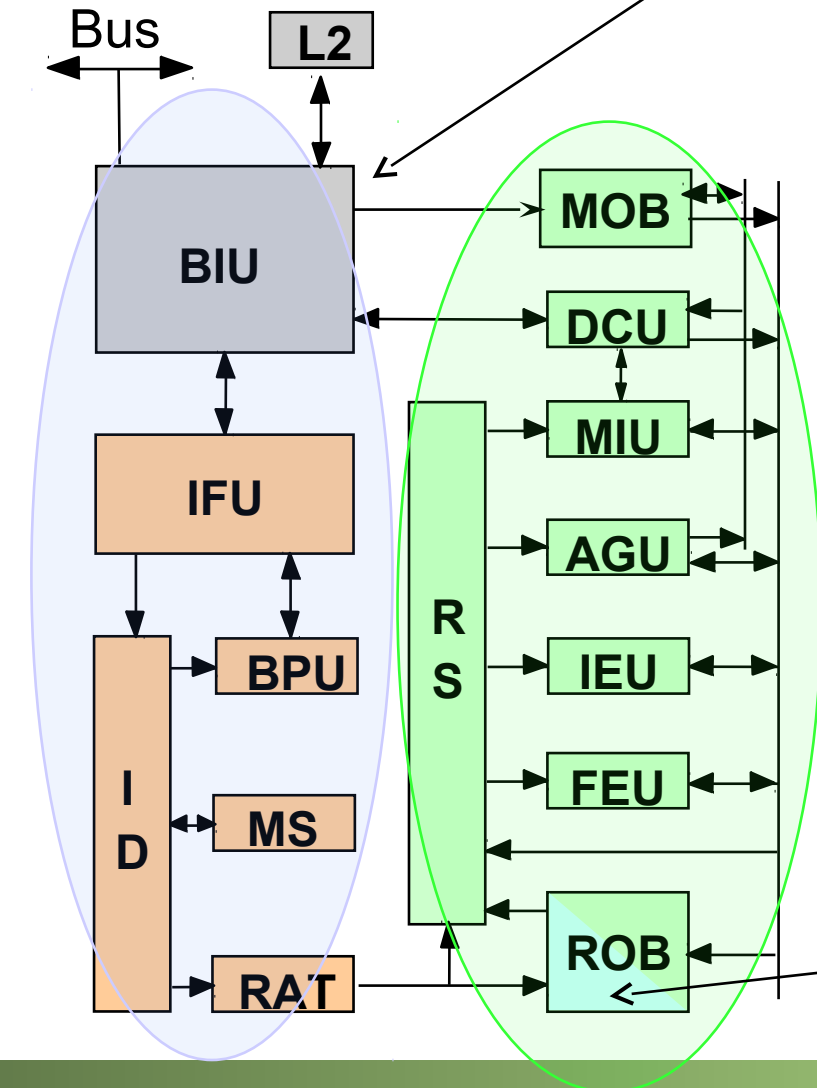
- Out of order, speculative, superscalar, renames registers

- **Retire**

- In order
- When speculation ends, commit
- Can simultaneously commit  $\leq 3$  (the “width” of the machine)

# P6 $\mu$ Arch

External



## • In-Order Front End

- BIU: Bus Interface Unit
- IFU: Instruction Fetch Unit (includes IC)
- BPU: Branch Prediction Unit
- ID: Instruction Decoder
- MS: Micro-Instruction Sequencer
- RAT: Register Alias Table

## • Out-of-order Core

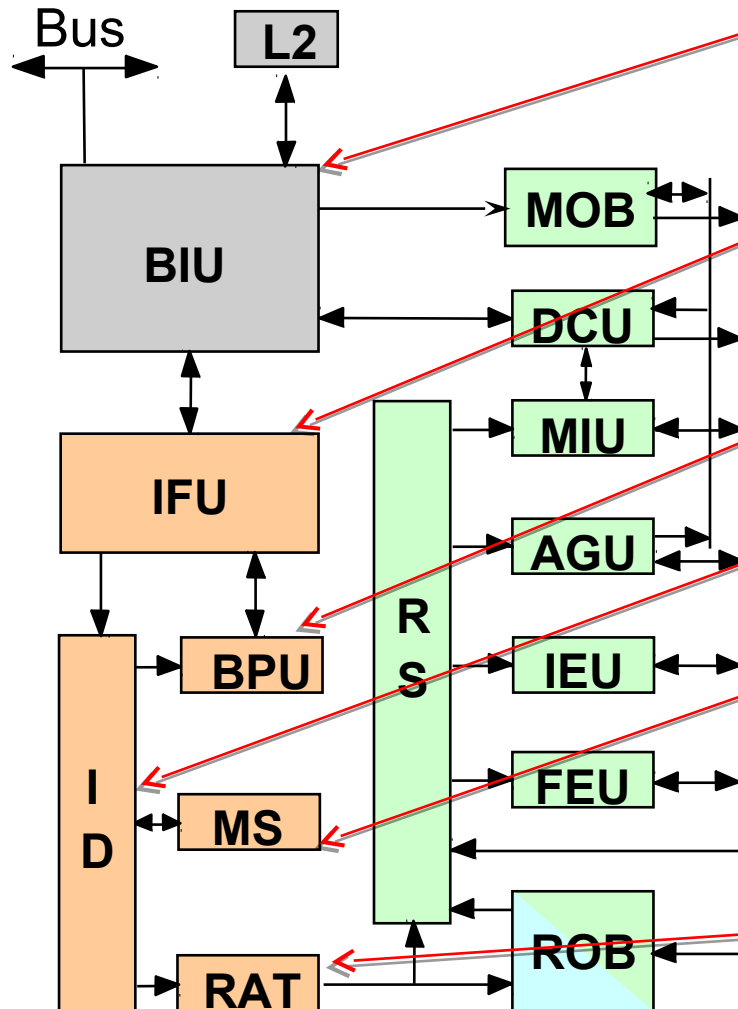
- ROB: Reorder Buffer
- RRF: Real Register File
- RS: Reservation Stations
- IEU: Integer Execution Unit
- FEU: Floating-point Execution Unit
- AGU: Address Generation Unit
- MIU: Memory Interface Unit
- DCU: Data Cache Unit
- MOB: Memory Order Buffer
- L2: Level 2 cache

## • In-Order Retire

# P6 $\mu$ Arch

## In-Order Front End

External



• **BIU:** Bus Interface Unit  
(fetches instructions)

• **IFU:** Instruction Fetch Unit  
(includes i-cache)

• **BPU:** Branch Prediction Unit

• **ID:** Instruction Decoder

• **MS:** Micro-Instruction Sequencer  
(complex ops are comprised of a  
sequence of  $\mu$ -ops; simple ops are  
comprised of only 1  $\mu$ -op)

• **RAT:** Register Alias Table  
(solves false dep.; most recent arch =>  
physical mapping)

# P6 $\mu$ Arch

## Out-of-order Core

- L2: Level 2 cache

- MOB: Memory Order Buffer

- DCU: Data Cache Unit

- MIU: Memory Interface Unit

- AGU: Address Generation Unit

- RRF: “Real” Register File  
(not shown; the machine’s state)

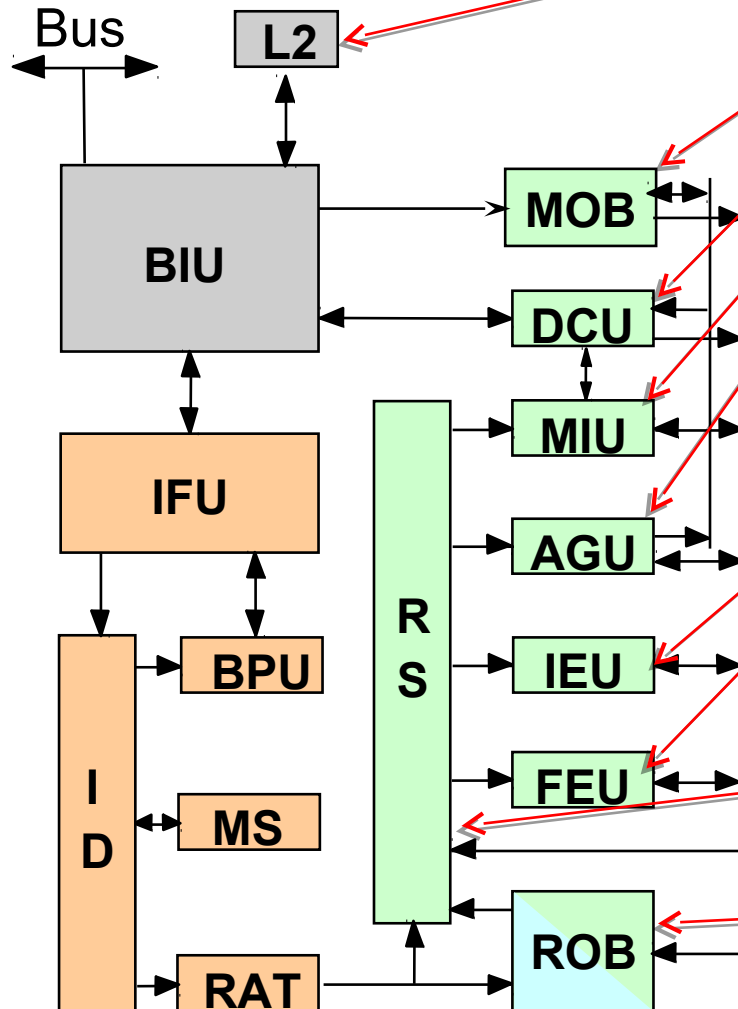
- IEU: Integer Execution Unit

- FEU: Floating-point Execution Unit


- RS: Reservation Stations  
(All those ops whose dependencies aren’t yet met; up to 20; when read, 5 ports to exe units)

- ROB: Reorder Buffer  
(The physical regs; one entry per op – the reg is the dest of the op; in order!)

External




# P6 Pipeline - 12 stages (10<=P6<=14)




**In-Order Front End**

- 1: Next IP
- 2: ICache lookup
- 3: ILD (instruction length decode)
- 4: rotate
- 5: ID1 (instruction decoded 1)
- 6: ID2 (instruction decoded 2)
- 7: RAT - rename sources & ALLOC - assign destinations
- 8: ROB - read sources  
RS - schedule data-ready uops for dispatch
- 9: RS - dispatch uops
- 10: EX
- 11-12: Retirement

**Out-of-order Core**

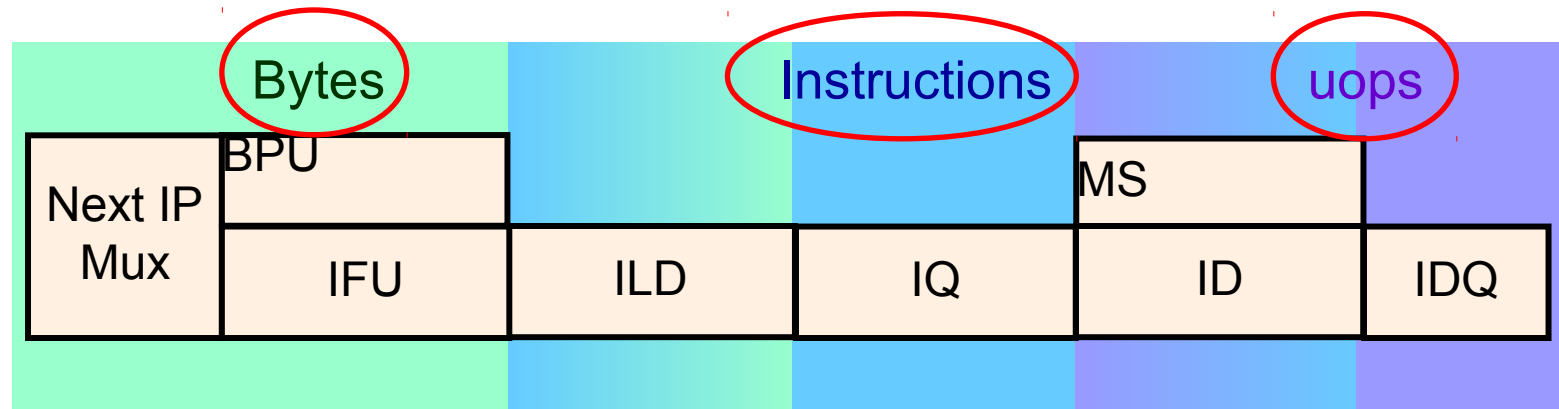


**In-order Retirement**





# In-Order Front End



- **BPU – Branch Prediction Unit – predict next fetch address**
- **IFU – Instruction Fetch Unit**
  - iTLB translates virtual to physical address
  - IC supplies 16byte/cyc (access L2 cache, maybe memory, on miss)
- **ILD – Induction Length Decode – split bytes to instructions**
- **IQ – Instruction Queue – buffer the instructions**
- **ID – Instruction Decode – decode instructions into uops**
- **MS – Micro-Sequencer – provides uops for complex instructions**
- **IDQ – Instruction Decode Queue – buffer the uops**

# Branch Prediction

- **Implementation**

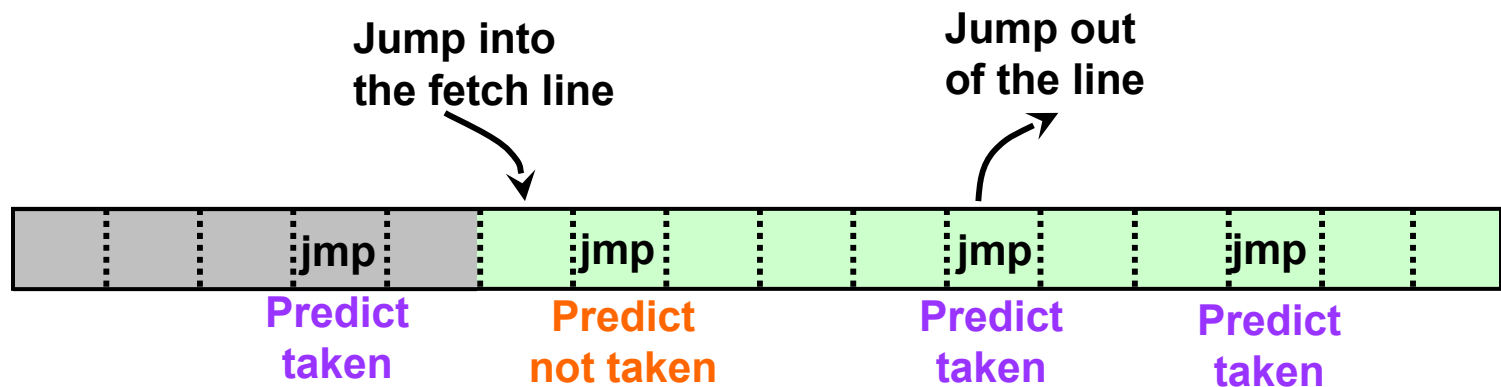
- Use local history to predict direction
- Need to predict multiple branches
- ⇒ Need to predict branches before previous branches are resolved
- ⇒ Branch history updated first based on prediction, later based on actual execution (speculative history)
- Target address taken from BTB

- **Prediction rate: ~92%**

- ~60 instructions between misprediction
- High prediction rate is very crucial for long pipelines
- Especially important for OOOE, speculative execution:
  - On misprediction all instructions following the branch in the instruction window are flushed
  - Effective size of the window is determined by prediction accuracy

# Branch Prediction – Clustering


- Given a fetched line (bytes), need to know which line to fetch next
  - Perhaps there's more than one branch in the line
  - We must use the 1<sup>st</sup> (leftmost) taken branch ( $\geq$  the fetched IP)



- Implemented by
  - Splitting IP into *set* + *tag* + *offset* (within line)
  - If there's a match
    - The offsets of the matching ways are ordered
    - Ways with offset smaller than the fetch IP offset are discarded
    - The 1st branch that's predicted taken is chosen as the predicted branch

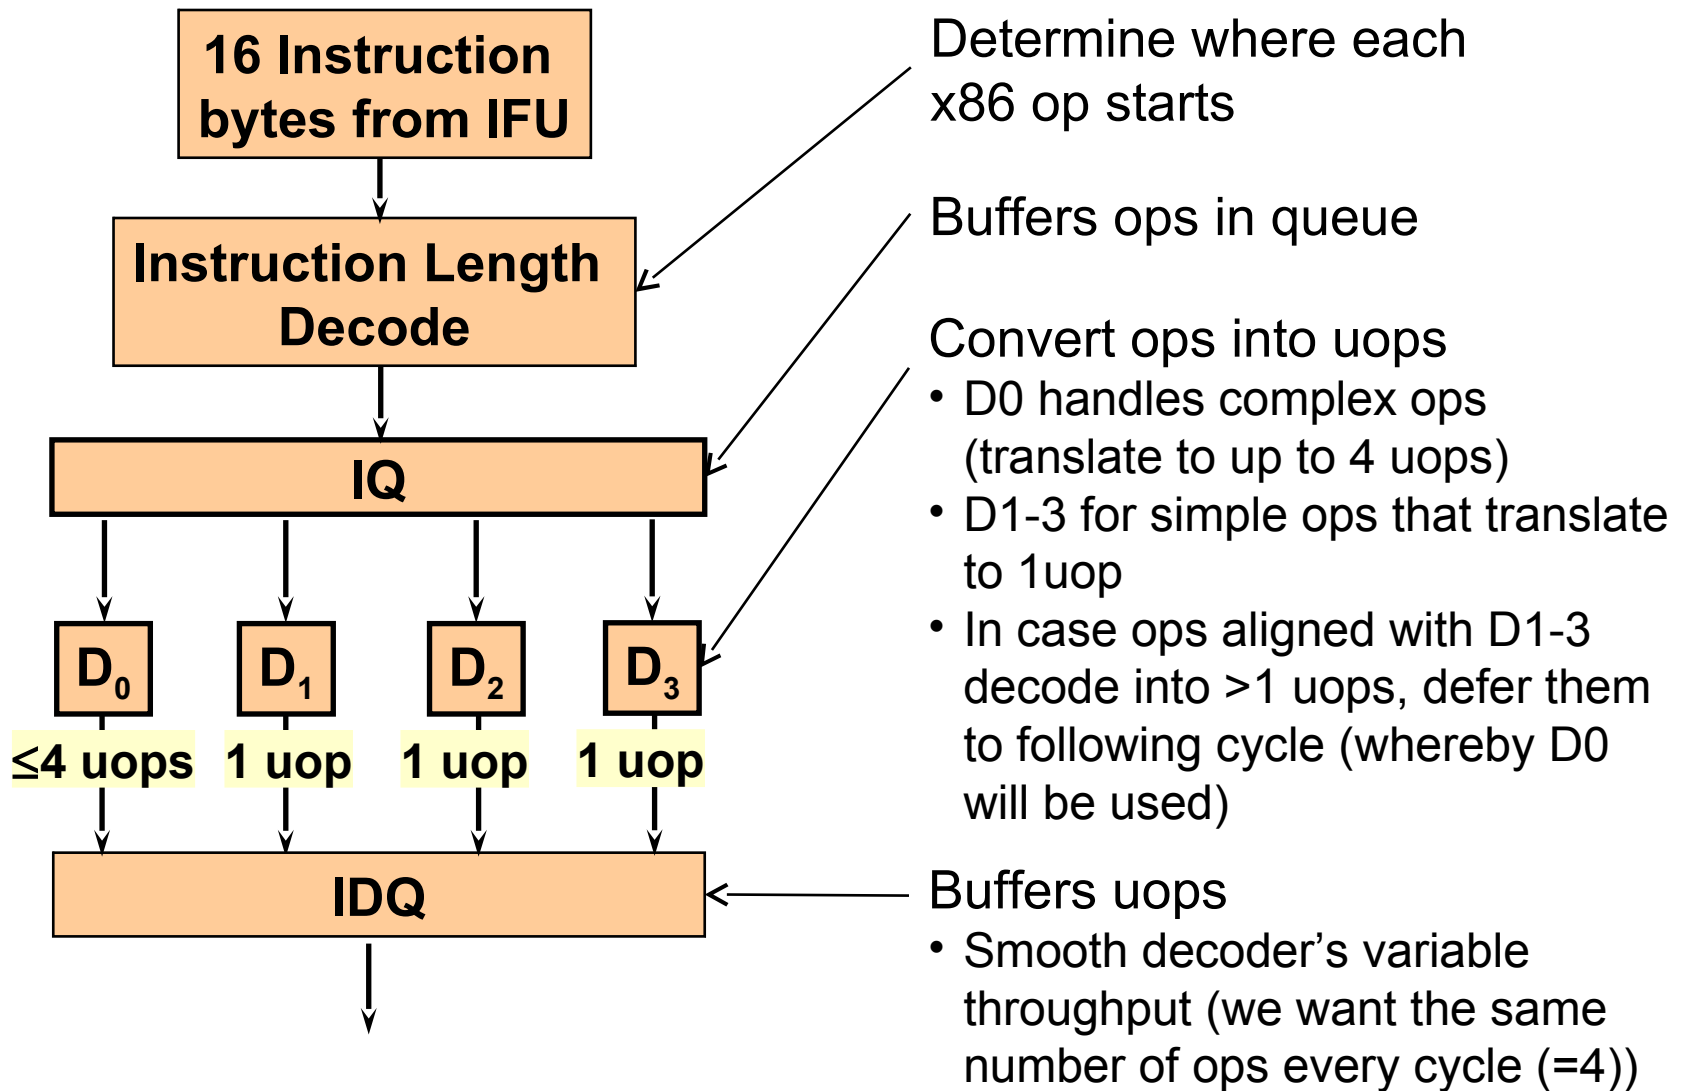
# The P6 BTB

- 2-level, local histories, per-set counters
- 4-way set associative: 512 entries in 128 sets



- Up to 4 branches can have a tag match

# In-Order Front End: Decoder



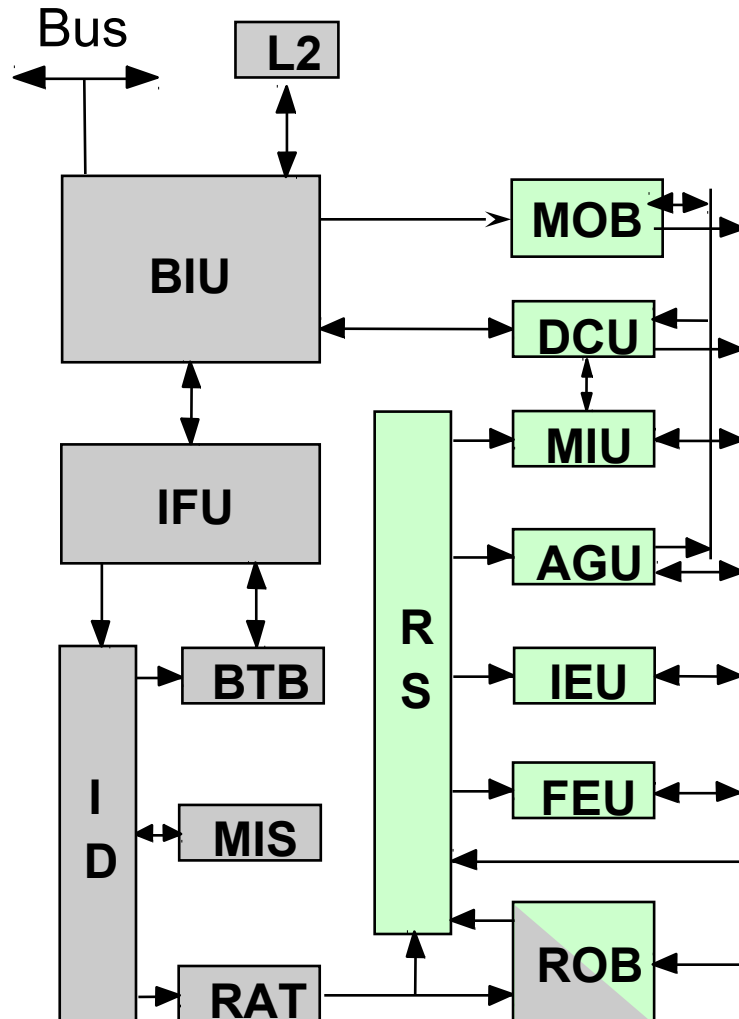
# Micro Operations (Uops)

- **Each CISC inst is broken into one or more RISC *uops***
  - Simplicity
    - Each uop is (relatively) simple
    - Canonical representation of src/dest (2 src, 1 dest)
  - But increased instruction count
- **Simple instructions translate to a few uops**
  - Typical uop count (it is not necessarily cycle count!)

Reg-Reg ALU/Mov inst:	1 uop
Mem-Reg Mov (load)	1 uop
Mem-Reg ALU (load + op)	2 uops
Reg-Mem Mov (store)	2 uops (st addr, st data)
Reg-Mem ALU (ld + op + st)	4 uops
- **Complex instructions need ucode**

# Out-of-order Core: ROB + RS

External



- **Reorder Buffer (ROB):**

- Holds “not yet retired” instructions
- 40 ordered entries (cyclic array)
- Retired in-order
- It’s possible some instruction already executed (their result known), but cannot be retired since
  - still have speculative status
  - and/or are waiting for previous instructions to retire in order

- **Reservation Stations (RS):**

- Holds “not yet executed” instructions
- 20 entries (subset of ROB)
- Up to 4 simultaneous ops can get in and out of RS simultaneously

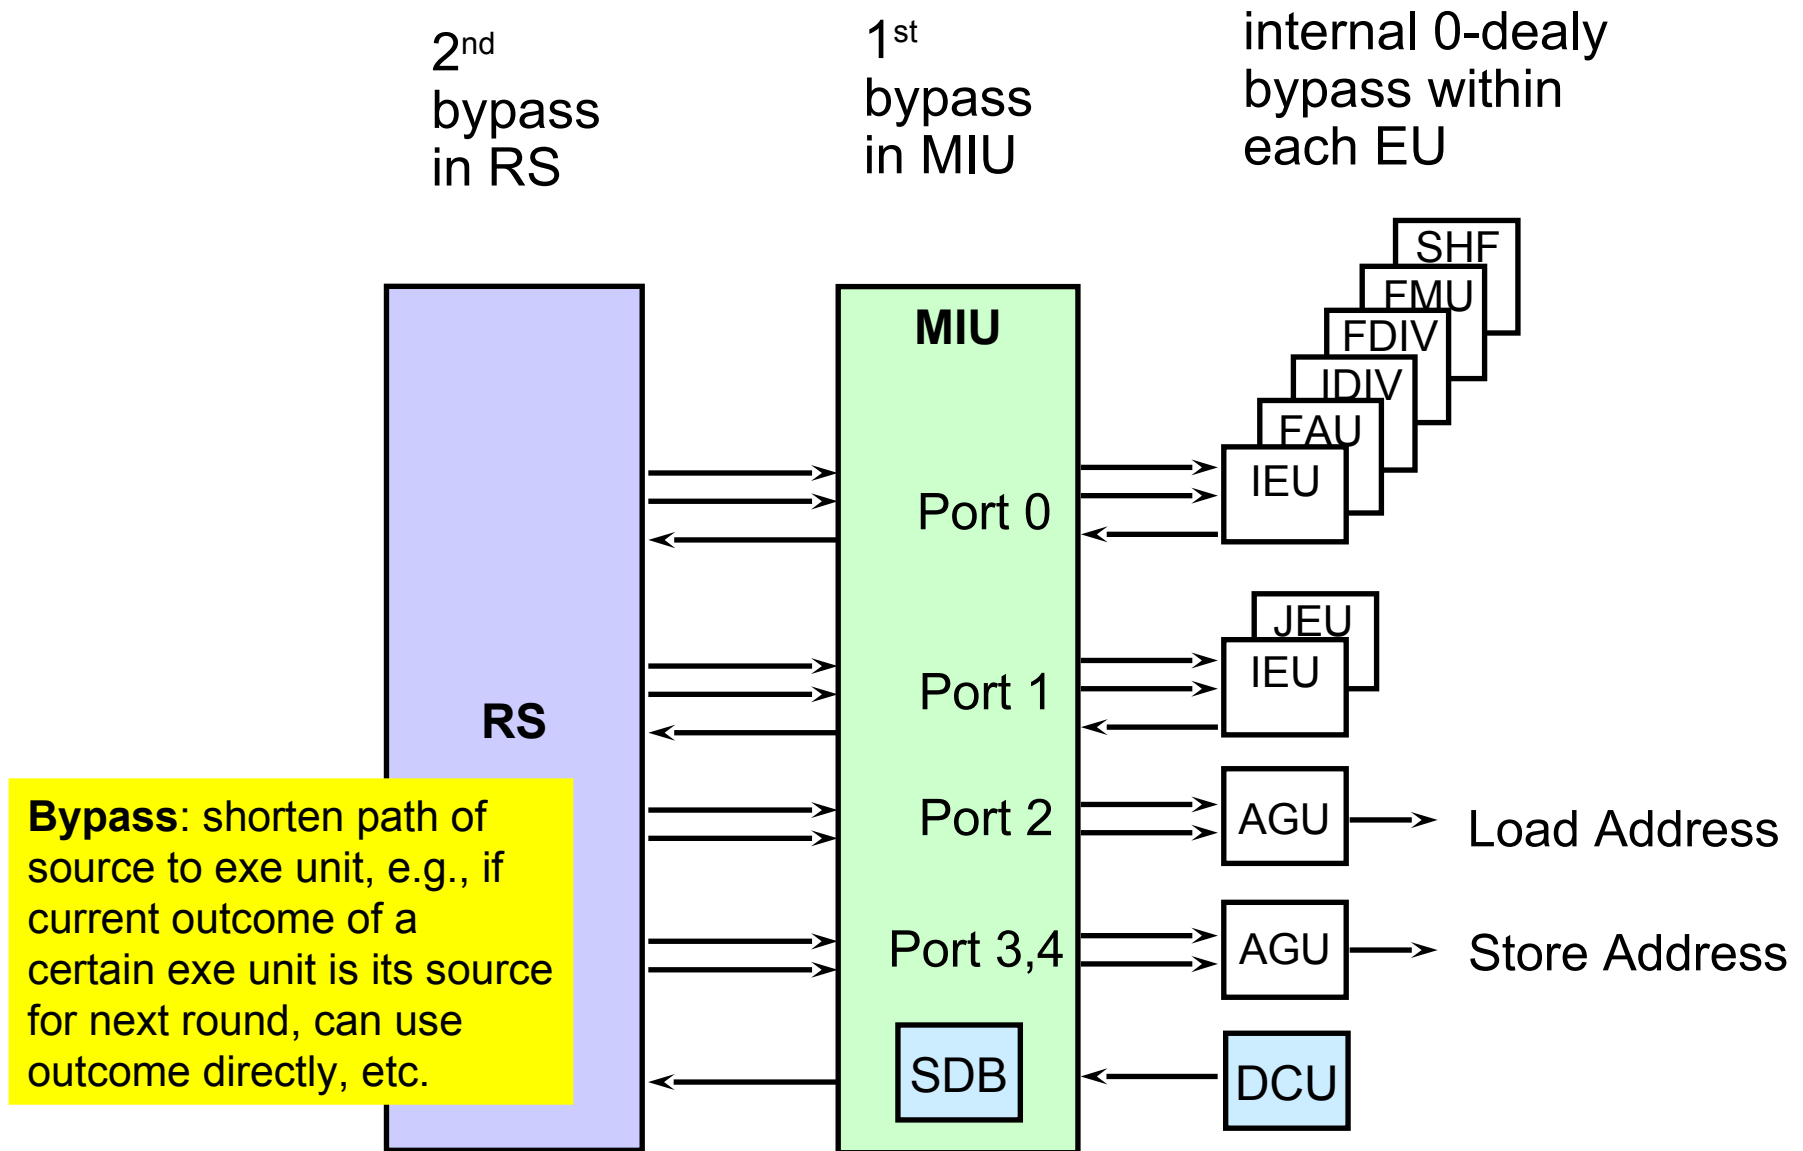
- **After execution**

- Results written to both ROB & possibly to RS (when source of other instructions)





# Out-of-order Core: execution units



# RAT & ALLOC

- There are  $\leq 4$  new uops/cyc; for each such uop
  - Perform register allocation & renaming, specifically...
- For each new uop , use RAT (**Register Alias Table**) to
  - **Source reg(s)**: map arch reg(s) to physical reg(s)
    - arch reg  $\Rightarrow$  latest phys reg that updated arch reg
  - **Target reg**: (1) allocate new phys reg; (2) update RAT accordingly
    - Now arch reg points to newly allocated phys reg (for next time)

RAT:

arch reg	phys reg#	location
EAX	0	RRF
EBX	19	ROB
ECX	23	ROB

- **The Allocator (Alloc)**
  - Assigns each uop with new ROB & RS entries
  - Write up the matching phys regs to RS (along with the rest of the uop)
  - Allocate Load & Store buffers in the MOB (for load & store ops)

# Reorder Buffer (ROB)

- Hold 40 uops which are “not yet committed”
  - Same order as program (cyclic array)
  - Provide large physical register space for reg renaming
  - A physical register is actually an item within a matching ROB entry
    - phys reg number = ROB entry number
    - phys reg = uop's target destination (there's always exactly one)
    - phys regs buffer the execution results until retirement

#entry	entryValid	dataValid	data (physical reg)	arch target reg
0	1	1	12H	EBX
1	1	1	33H	ECX
2	1	0	xxx	ESI
...				
39	0	0	xxx	XXX

- *Valid data* is set after uop executed (& result written to physical reg)

# RRF – Real Register File

- **Holds the Architectural Register File**
  - Architectural Register are numbered: 0 – EAX, 1 – EBX, ...
  - This is “the state” of the chip (can’t roll back)
- **The value of an architectural register**
  - Is the **value written to it by the last committed uop** (which writes to that reg)
  - So long as we don’t change the RRF, we don’t change the state

**RRF:**

#entry	Arch Reg Data
0 (EAX)	9AH
1 (EBX)	F34H

# Uop flow through the ROB

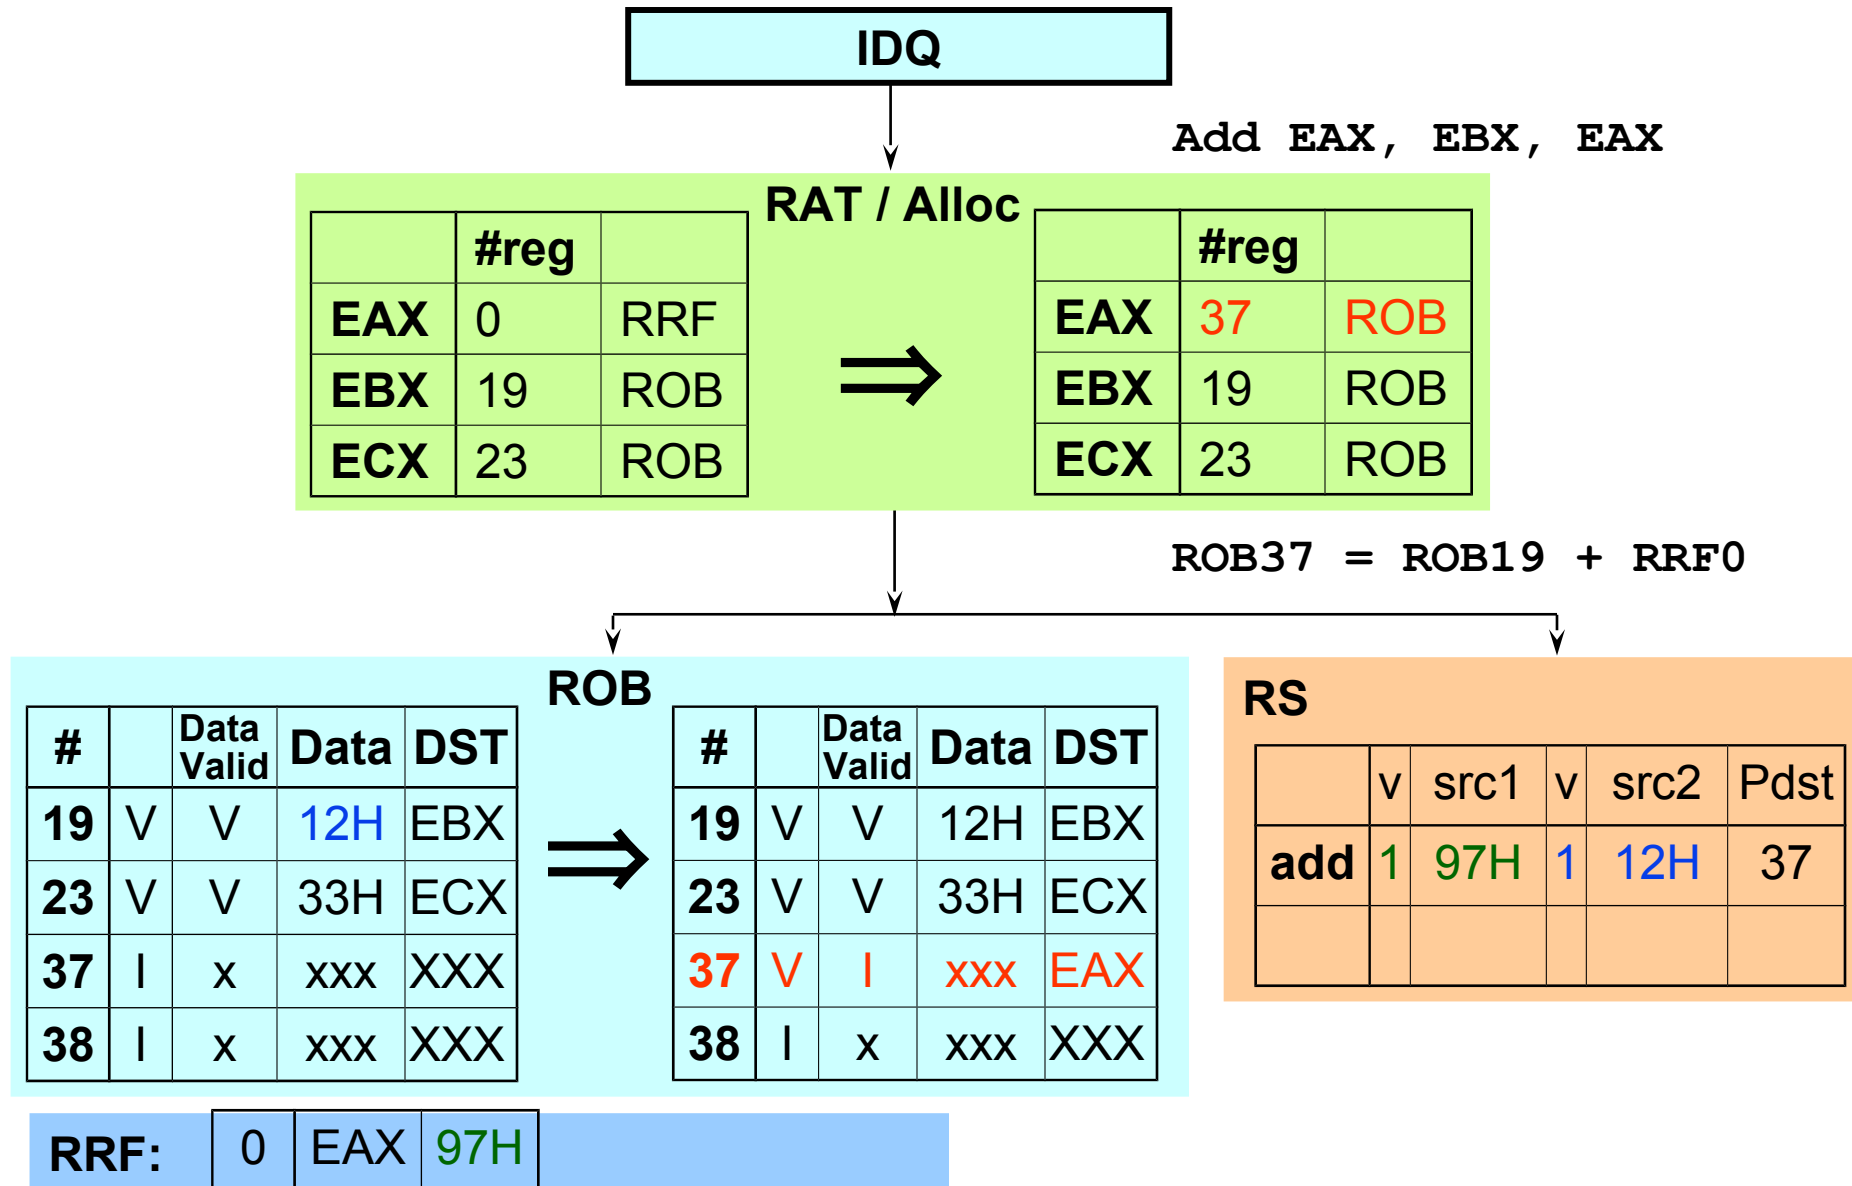
- **Uops are entered in order**
  - Registers renamed by the entry # (there's a head and a tail)
- **Once assigned**
  - Execution order unimportant
- **After execution:**
  - Entries marked “*executed*” (*dataValid=1*) & wait for retirement
  - Retirement occurs once all prior instruction have retired
  - $\Rightarrow$  Commit architectural state only after speculation was resolved
- **Retirement**
  - Detect exceptions and misprediction
    - Branch result might impact uops down the road
    - Initiate repair to get machine back on track
  - Update “*real*” regs (in RRF) with value of renamed (phys) regs
  - Update memory
  - Clear ROB entry

...

# Reservation station (RS)

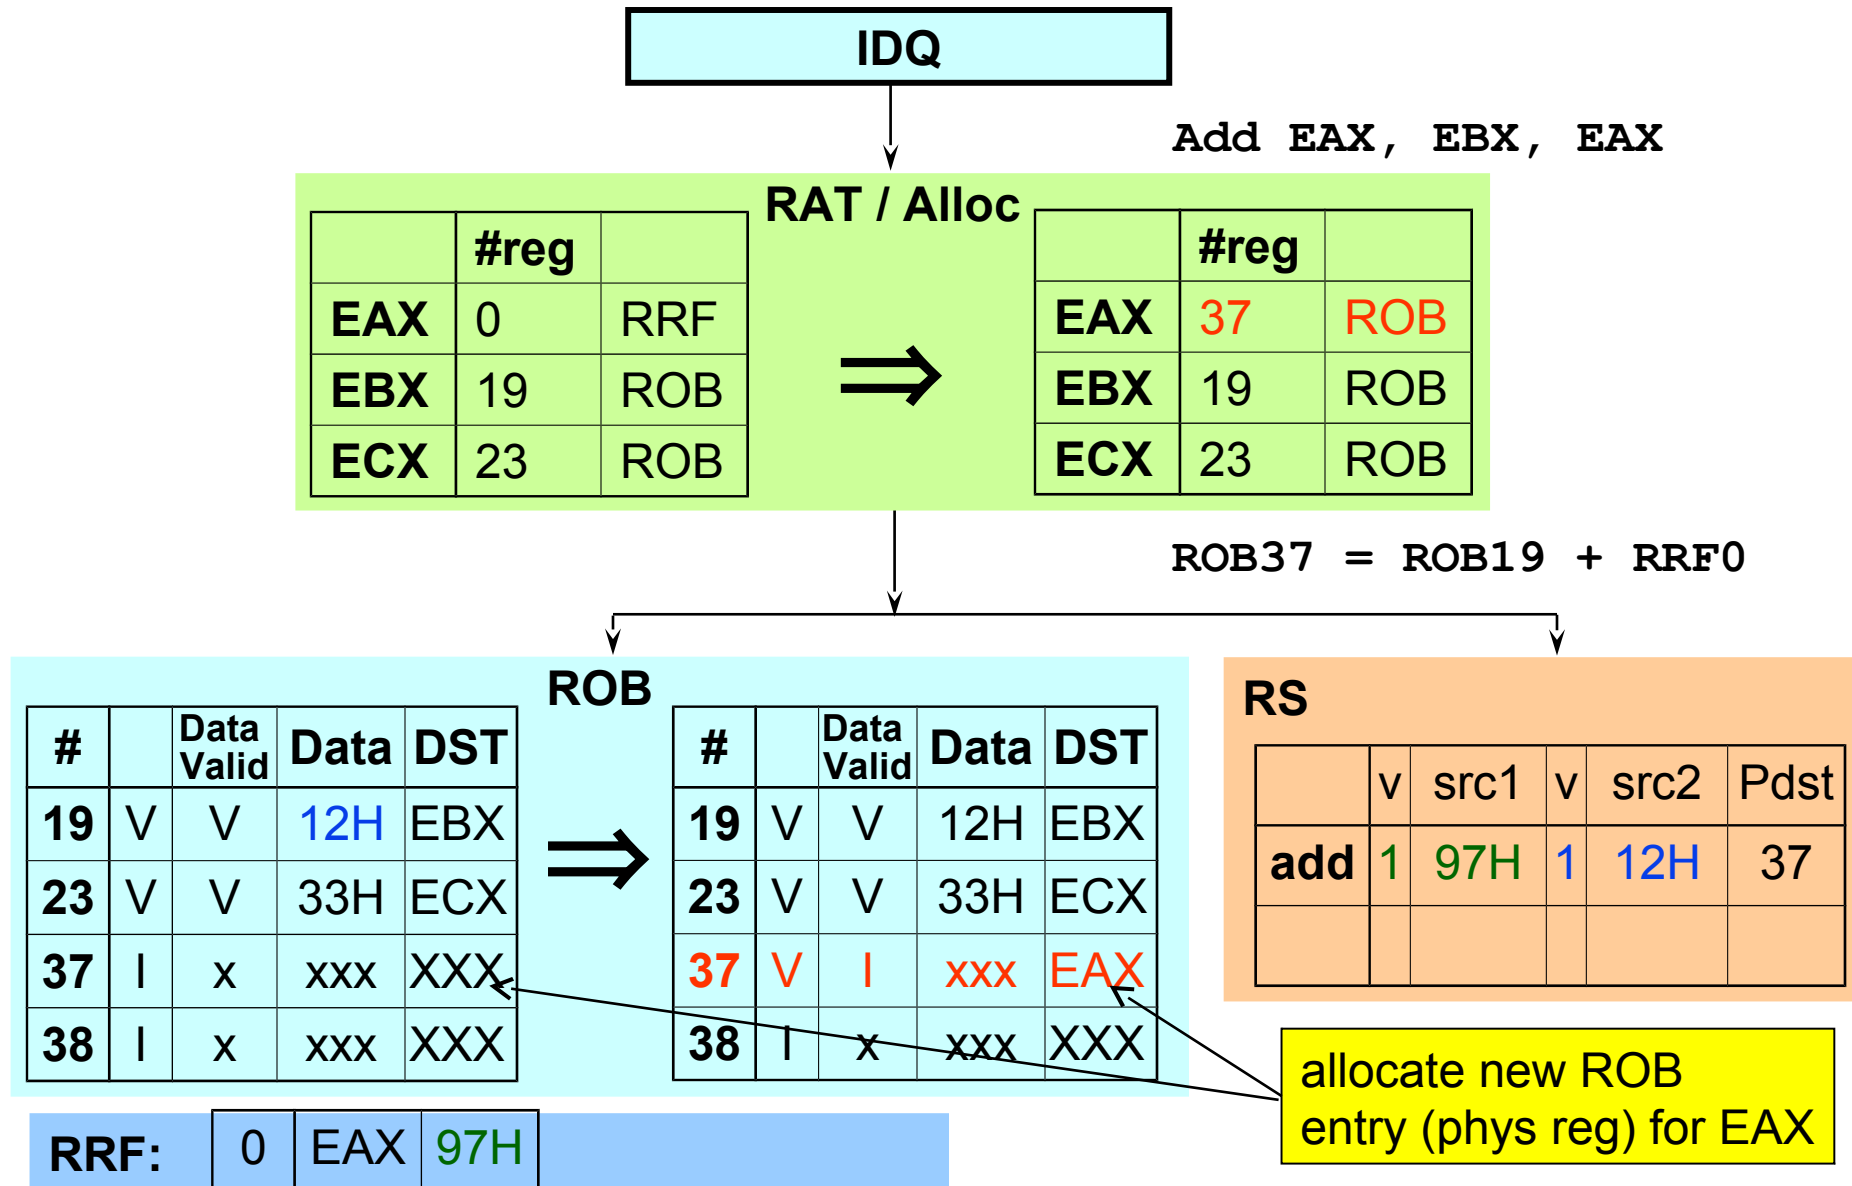
- **Pool of all “not yet executed” uops**
  - Holds the uop code & source data (until it is dispatched=scheduled)
- **When a uop is allocated in RS, operand values are updated**
  - If operand is arch reg => value taken from the RRF
  - If operand is phys reg (with *dataValid* =1) => value taken from ROB
  - If operand is phys reg (with *dataValid*=0) => wait for value
- **The RS maintains operands status “*ready* / *not-ready*”**
  - Each cycle, executed uops make more operands “*ready*”
    - RS arbitrates WB busses between exe units
    - RS monitors WB bus to capture data needed by waiting uops
    - Data can bypass directly from WB bus to exe unit (like we’ve seen)
  - Uops whose all operands are *ready* can be *dispatched*
    - Dispatcher chooses which *ready* uop to execute next
    - Dispatcher sends chosen uops to appropriate functional units
    - Need appropriate functional unit to be vacant

# Register Renaming example

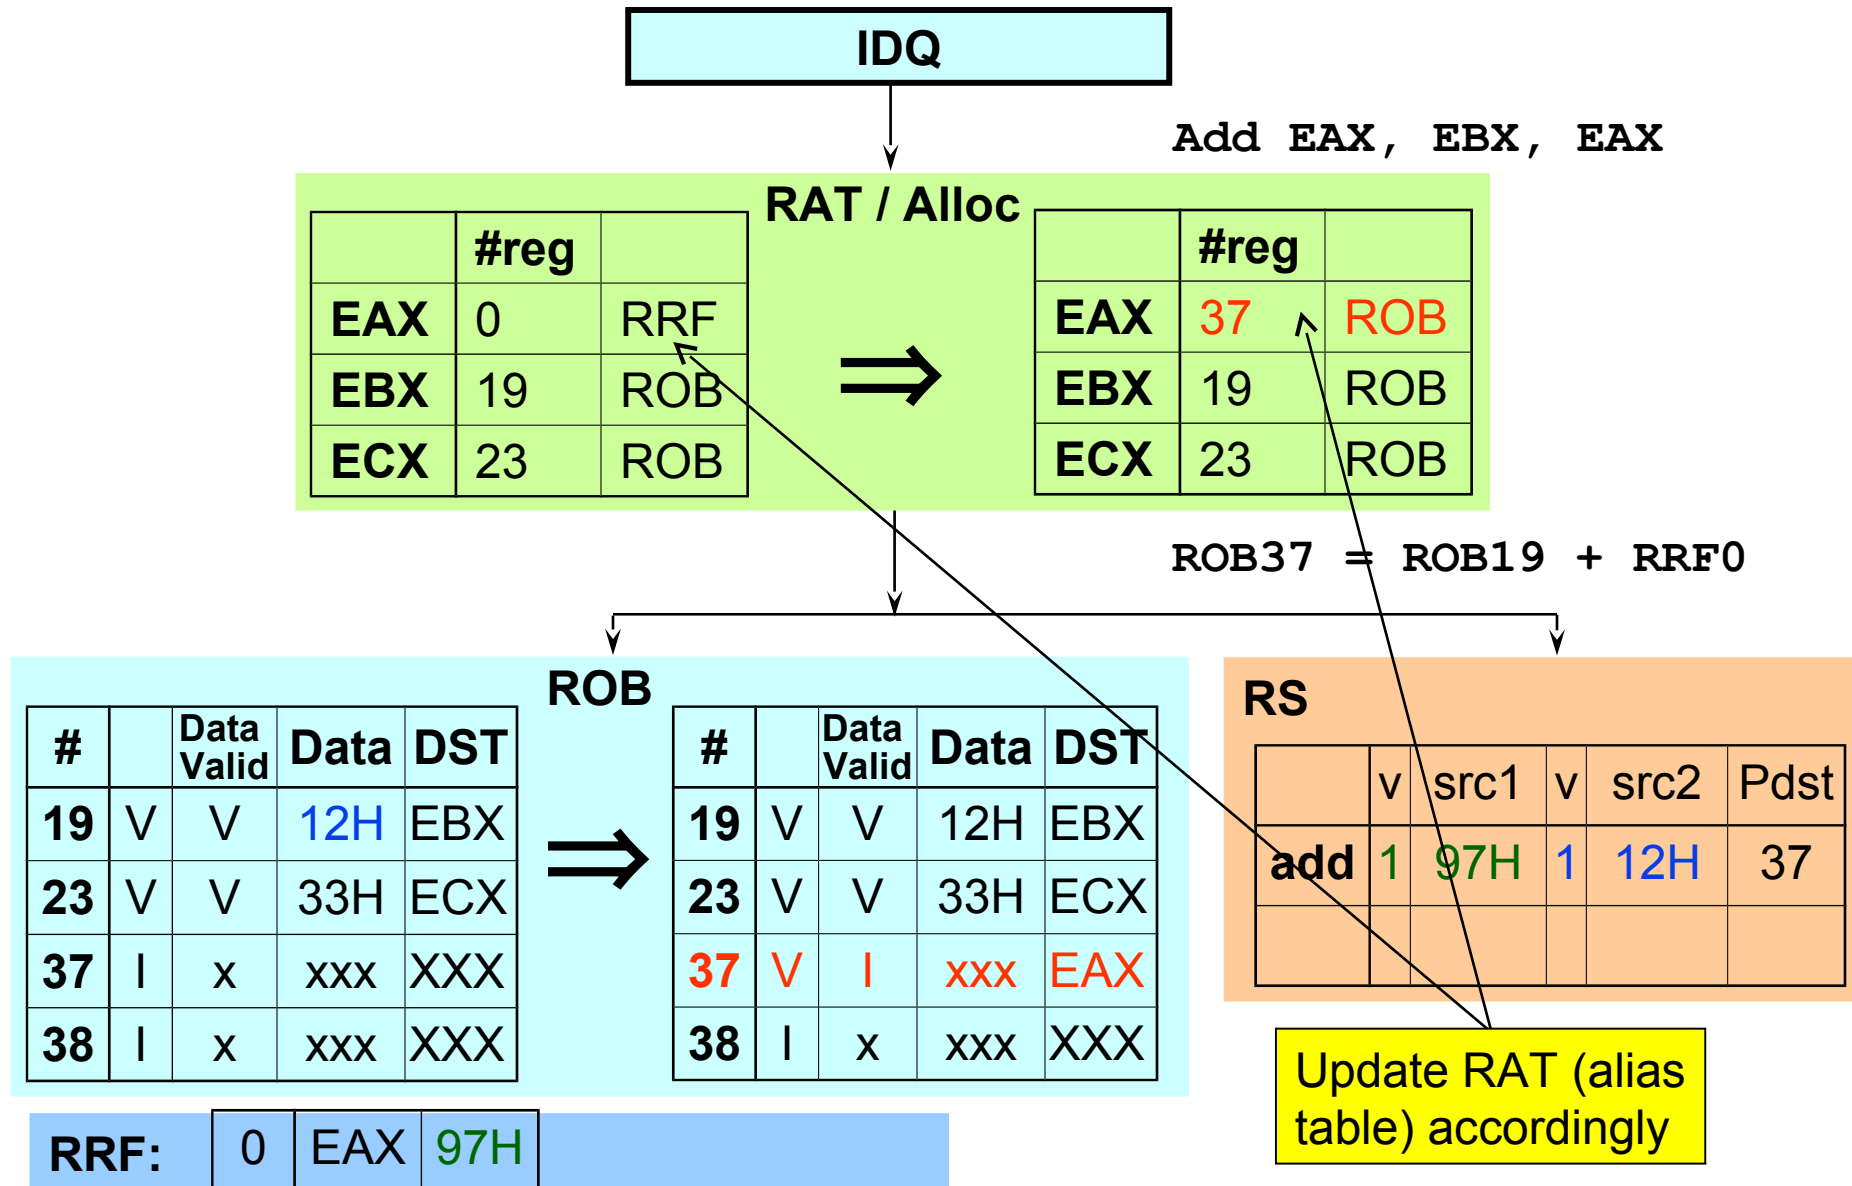




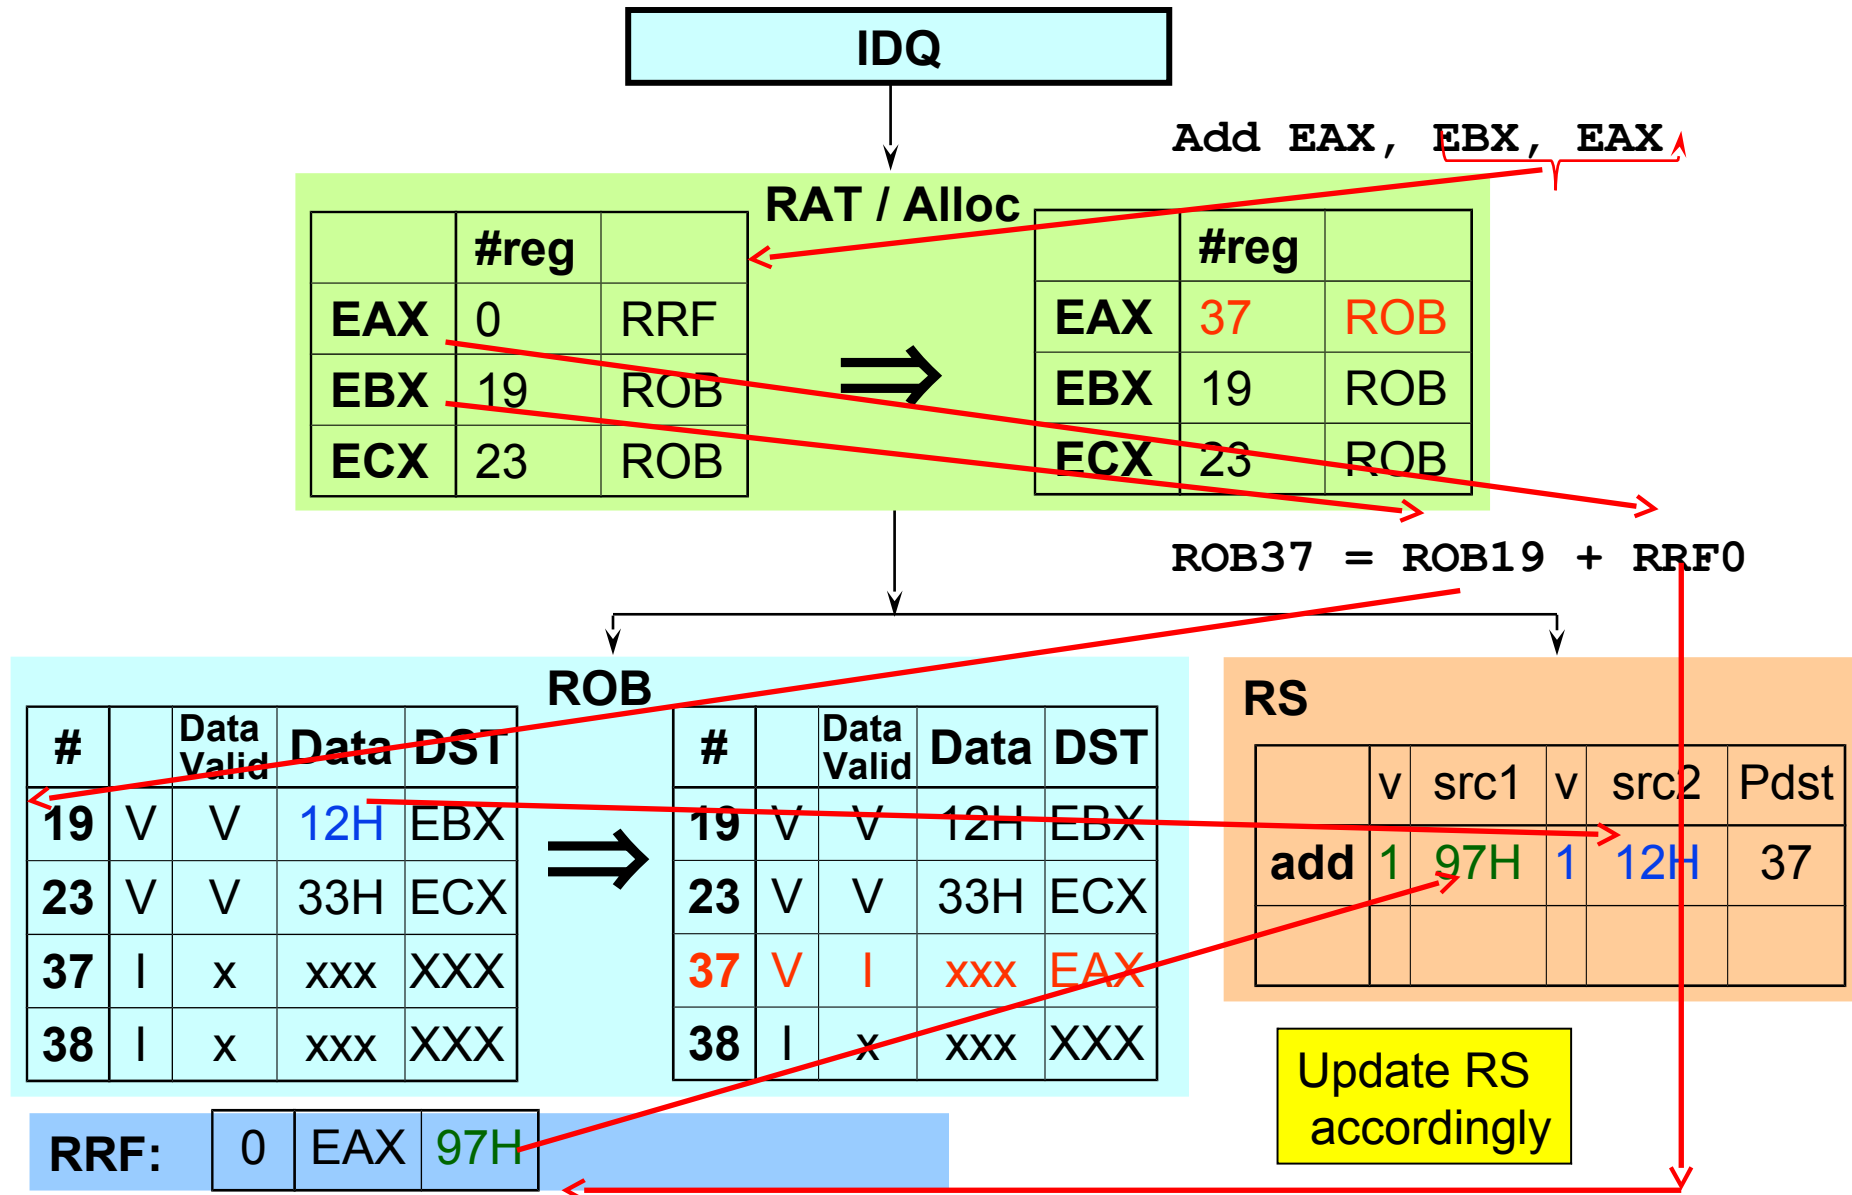
# Register Renaming example



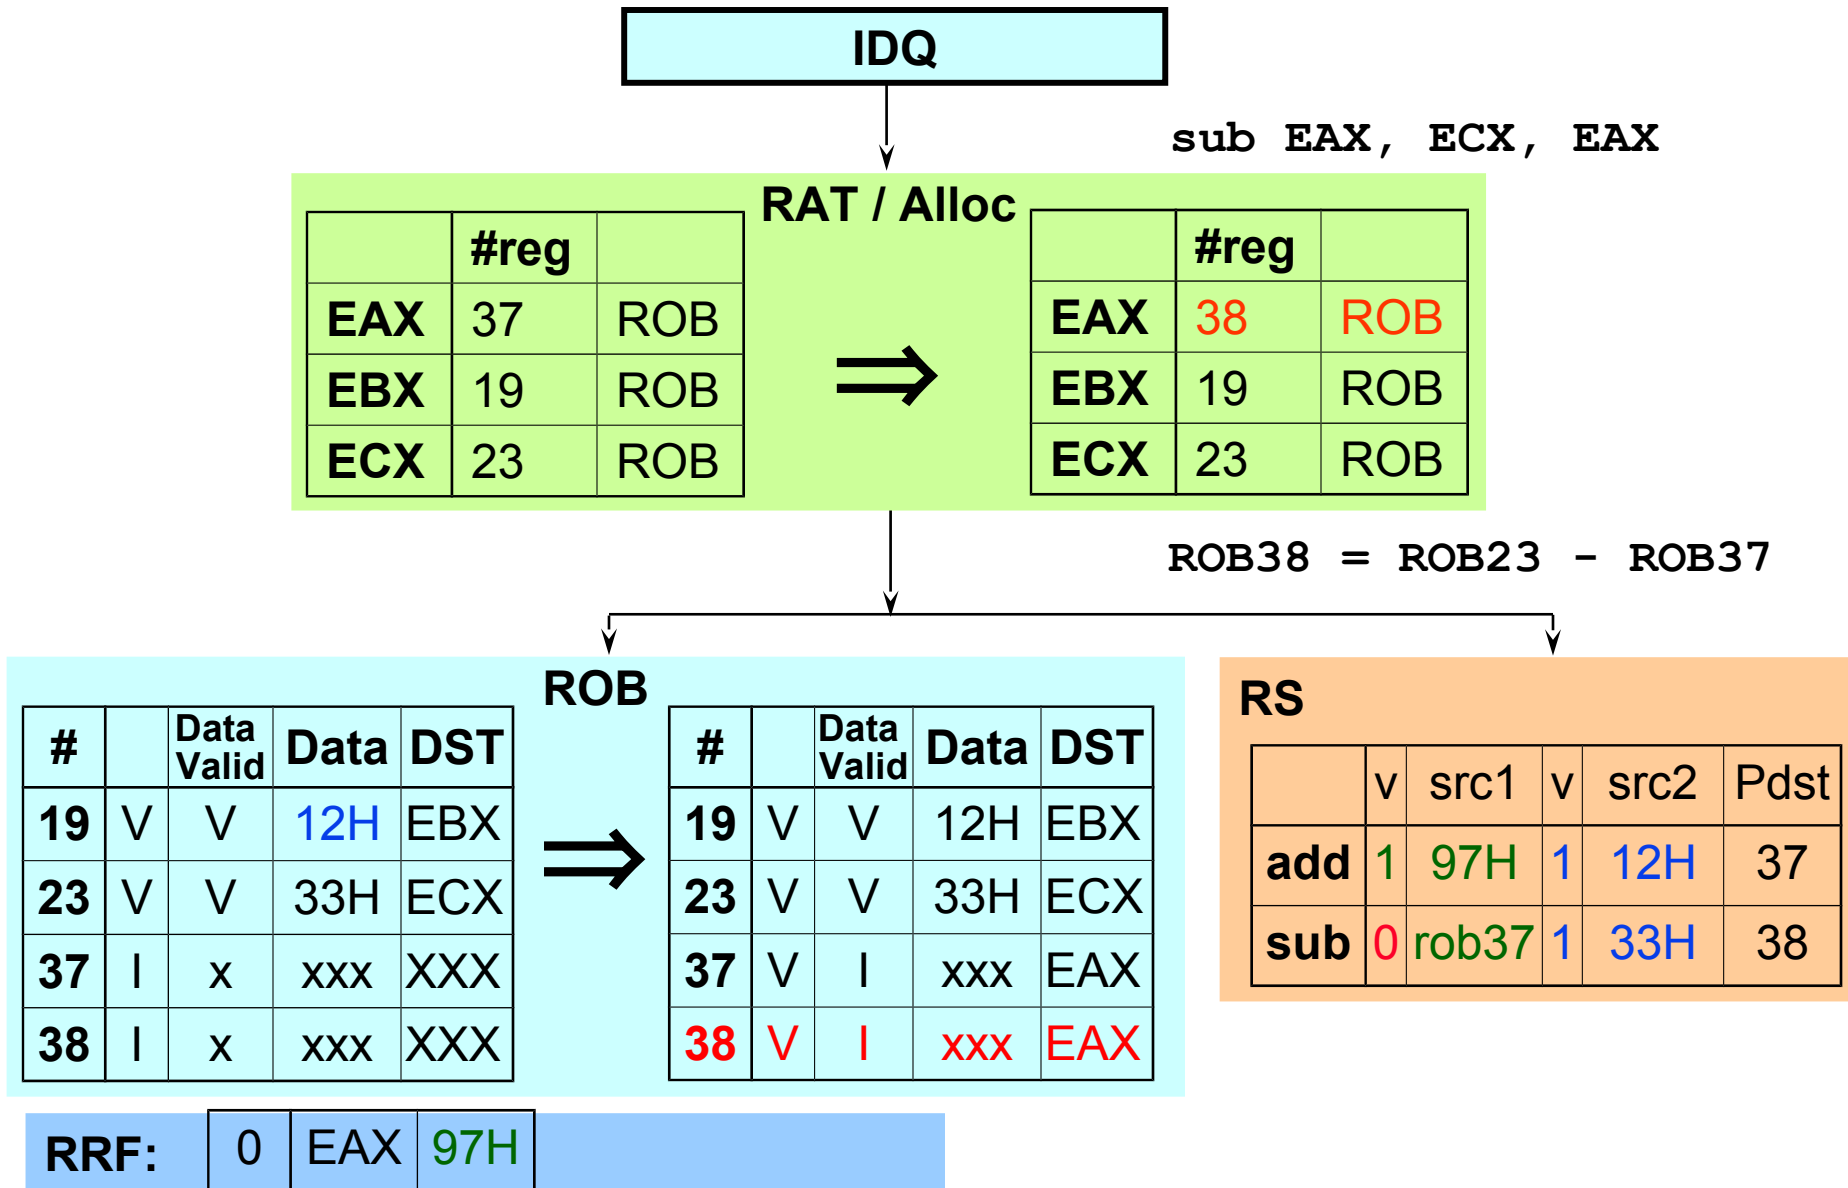
# Register Renaming example



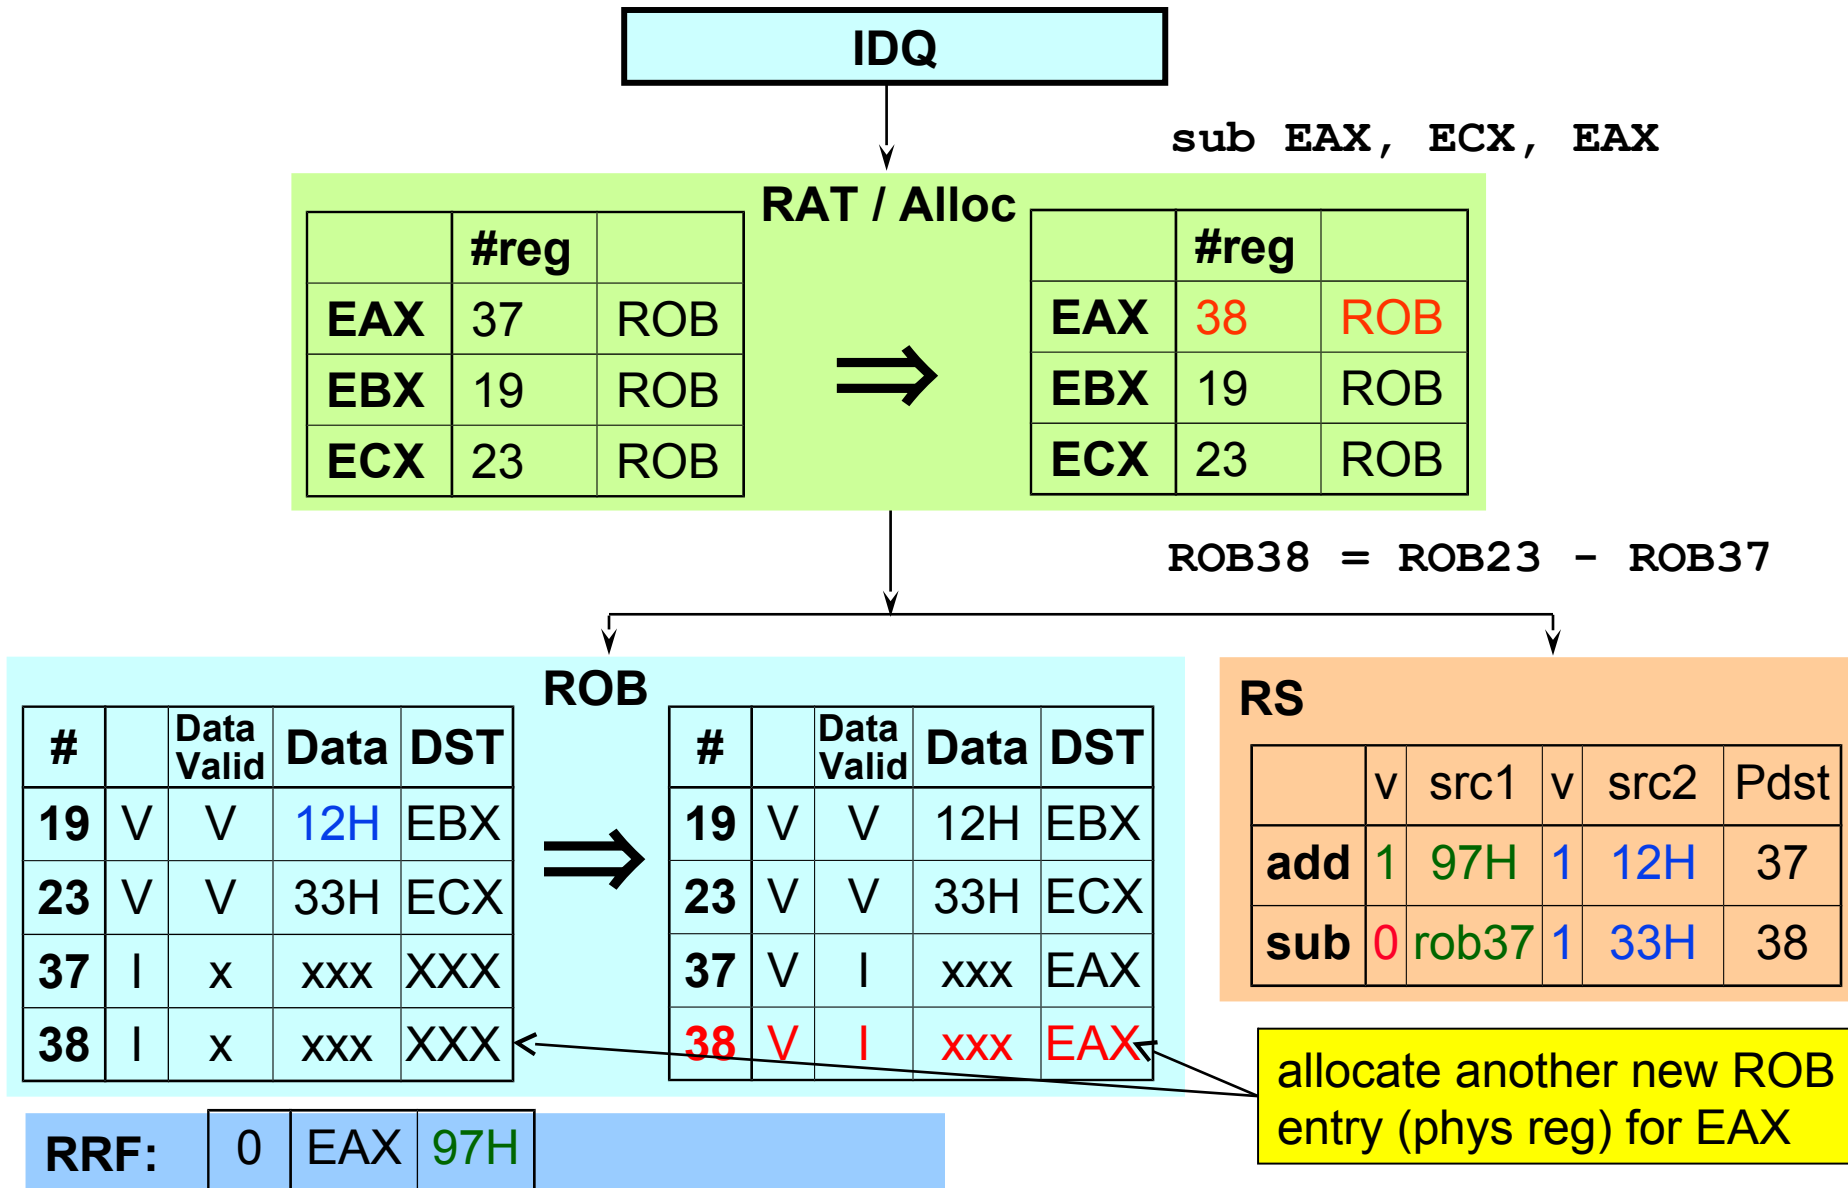
# Register Renaming example



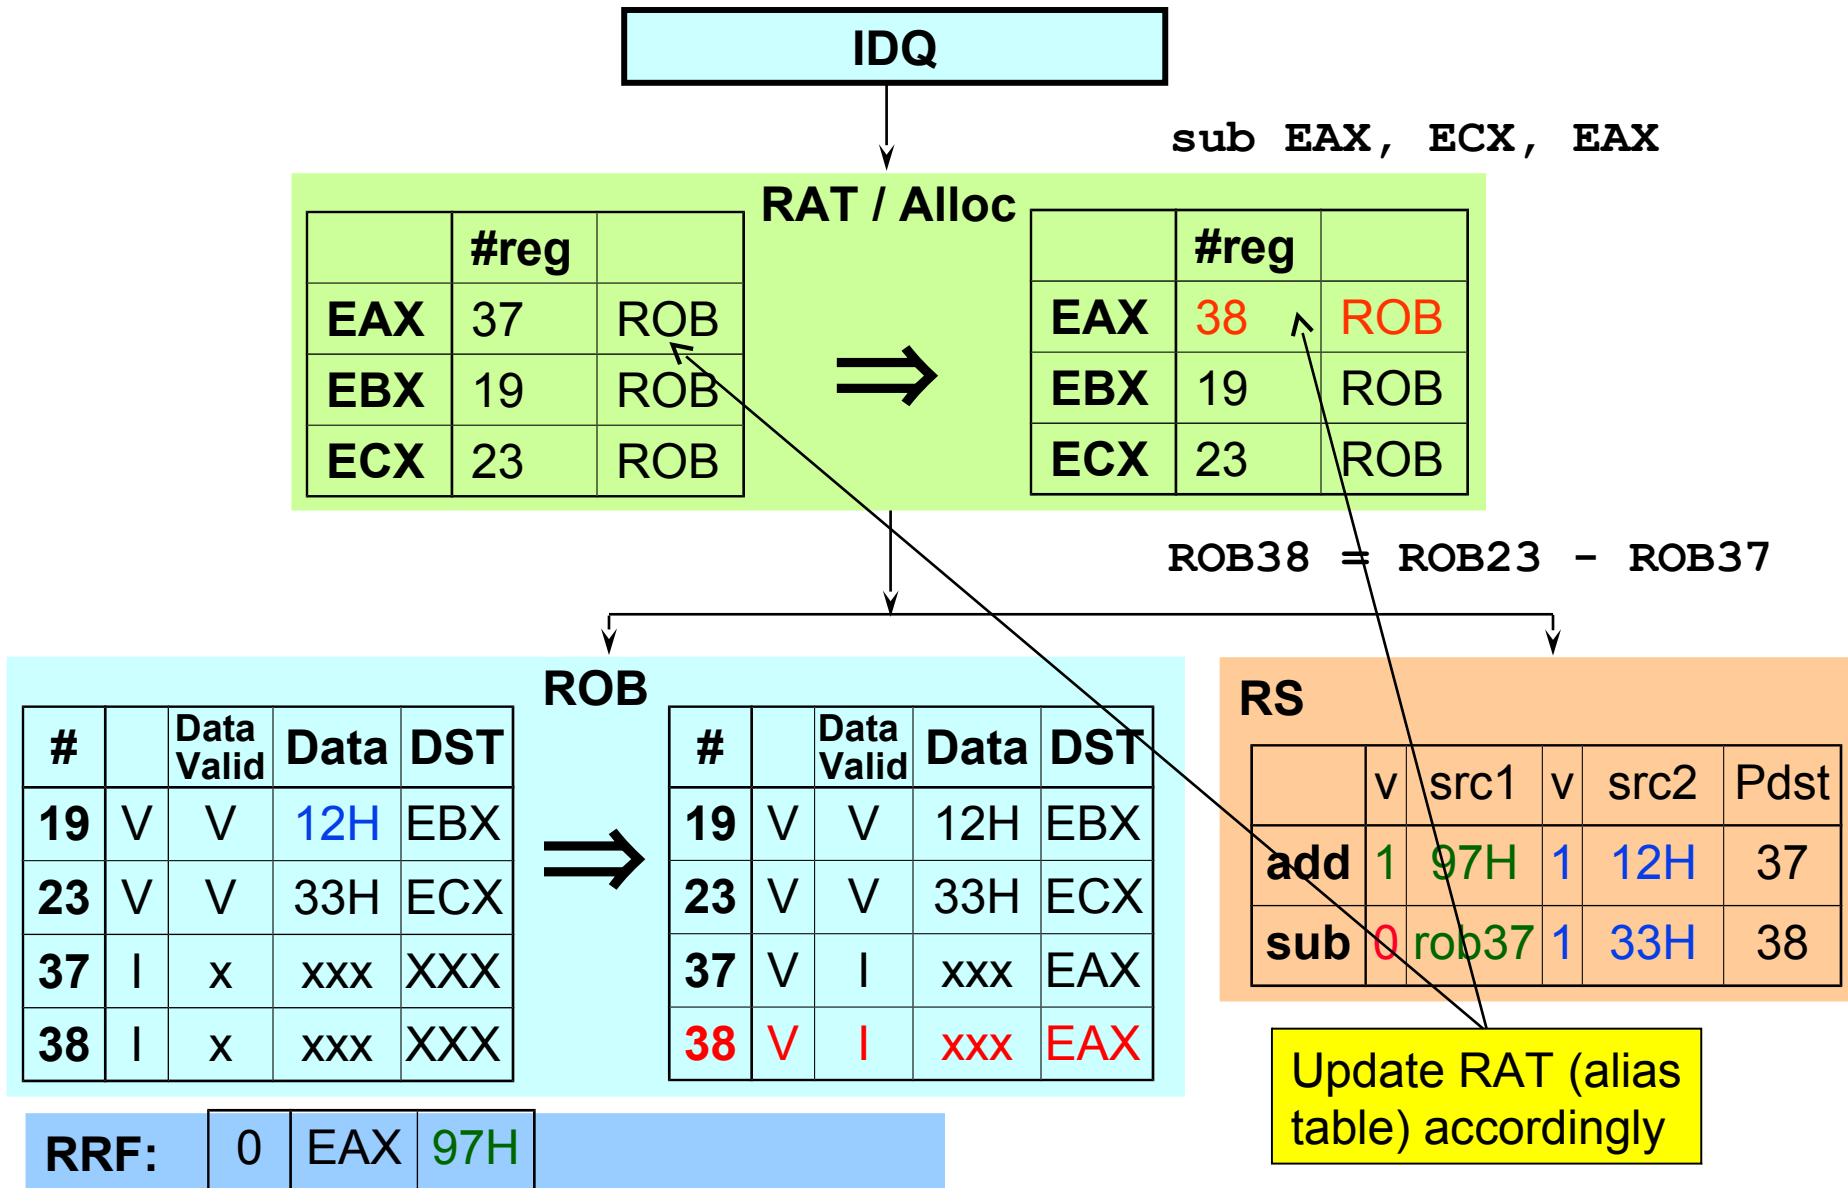
# Register Renaming example (2)



# Register Renaming example (2)



# Register Renaming example (2)



# Register Renaming example (2)

IDQ

sub EAX, ECX, EAX

RAT / Alloc		
	#reg	
EAX	37	ROB
EBX	19	ROB
ECX	23	ROB

⇒

	#reg	
EAX	38	ROB
EBX	19	ROB
ECX	23	ROB

ROB38 = ROB23 - ROB37

ROB				
#		Data Valid	Data	DST
19	V	V	12H	EBX
23	V	V	33H	ECX
37	I	x	xxx	XXX
38	I	x	xxx	XXX

⇒

#		Data Valid	Data	DST
19	V	V	12H	EBX
23	V	V	33H	ECX
37	V	I	xxx	EAX
38	V	I	xxx	EAX

RS					
	v	src1	v	src2	Pdst
add	1	97H	1	12H	37
sub	0	rob37	1	33H	38

RRF:	0	EAX	97H
------	---	-----	-----