Lecture 9 Directory Based Multiprocessors

Slides were used during lectures by David Patterson, Berkeley, spring 2006

Review

- Caches contain all information on state of cached memory blocks
- Snooping cache over shared medium for smaller MP by invalidating other cached copies on write
- Sharing cached data ⇒ Coherence (values returned by a read), Consistency (when a written value will be returned by a read)

Outline

- Review
- · Directory-based protocols and examples
- Synchronization
- Consistency
- · Cross Cutting Issues
- · Fallacies and Pitfalls
- Sun T1 ("Niagara") Multiprocessor
- · Microprocessor Comparison
- Conclusion

A Cache Coherent System Must:

- Provide set of states, state transition diagram, and actions
- Manage coherence protocol
 - (0) Determine when to invoke coherence protocol
 - (a) Find info about state of block in other caches to determine action
 - » whether need to communicate with other cached copies
 - (b) Locate the other copies
 - (c) Communicate with those copies (invalidate/update)
- (0) is done the same way on all systems
 - state of the line is maintained in the cache
- protocol is invoked if an "access fault" occurs on the line
- · Different approaches distinguished by (a) to (c)

Bus-based Coherence

- · All of (a), (b), (c) done through broadcast on bus
 - faulting processor sends out a "search"
 - others respond to the search probe and take necessary
- · Could do it in scalable network too
 - broadcast to all processors, and let them respond
- Conceptually simple, but broadcast doesn't scale with p
 - on bus, bus bandwidth doesn't scale
 - on scalable network, every fault leads to at least p network transactions
- · Scalable coherence:
 - can have same cache states and state transition diagram
 - different mechanisms to manage protocol

Scalable Approach: Directories

- Every memory block has associated directory information
 - keeps track of copies of cached blocks and their states
 - on a miss, find directory entry, look it up, and communicate only with the nodes that have copies if necessary
 - in scalable networks, communication with directory and copies is through network transactions
- Many alternatives for organizing directory information

Basic Operation of Directory k processors · With each cache-block in memory: k presence-bits, 1 dirty-bit • With each cache-block in cache: 1 valid bit, and 1 dirty (owner) bit

• Read from main memory by processor i:

19111911111

- If dirty-bit OFF then { read from main memory; turn p[i] ON; }
- if dirty-bit ON then { recall line from dirty proc (cache state to shared); update memory; turn dirty-bit OFF; turn p[i] ON; supply recalled data to i;}
- Write to main memory by processor i:
 - If dirty-bit OFF then { supply data to i; send invalidations to all caches that have the block; turn dirty-bit ON; turn p[i] ON; ... }
 - ...

Directory Protocol

- Similar to Snooping Protocol: Three states
 - Shared: ≥ 1 processors have data, memory up-to-date
 - Uncached (no processor has it; not valid in any cache)
 - Exclusive: 1 processor (owner) has data; memory out-of-date
- · In addition to cache state, must track which processors have data when in the shared state (usually bit vector, 1 if processor has copy)
- · Keep it simple(r):
 - Writes to non-exclusive data
 ⇒ write miss
 - Processor blocks until access completes
 - Assume messages received and acted upon in order sent

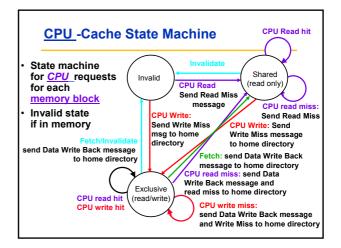
Directory Protocol

- · No bus and don't want to broadcast:
 - interconnect no longer single arbitration point
 - all messages have explicit responses
- · Terms: typically 3 processors involved
 - Local node where a request originates
 - Home node where the memory location of an address resides
- Remote node has a copy of a cache block, whether exclusive or shared
- Example messages on next slide: P = processor number, A = address

Directory Protocol Messages (Fig 4.22) Source Destination Msg Content Message type Read miss Local cache P. A Processor P reads data at address A; make P a read sharer and request data Write miss Local cache P, A Processor P has a write miss at address A; make P the exclusive owner and request data Invalidate a shared copy at address A Fetch the block at address A and send it to its home directory; change the state of A in the remote cache to shared Fetch the block at address A and send it to its home directory; invalidate the block in the cache Data value reply Return a data value from the home memory (read miss response) - Write back a data value for address A (invalidate response)

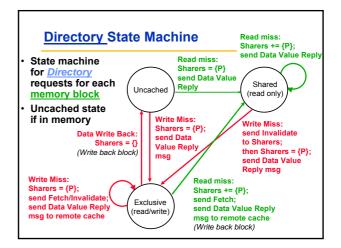
State Transition Diagram for One Cache Block in Directory Based System

- States identical to snooping case; transactions very similar
- Transitions caused by read misses, write misses, invalidates, data fetch requests
- Generates read miss & write miss message to home directory
- Write misses that were broadcast on the bus for snooping ⇒ explicit invalidate & data fetch requests
- Note: on a write, a cache block is bigger, so need to read the full cache block



State Transition Diagram for Directory

- Same states & structure as the transition diagram for an individual cache
- 2 actions: update of directory state & send messages to satisfy requests
- · Tracks all copies of memory block
- Also indicates an action that updates the sharing set, Sharers, as well as sending a message



Example Directory Protocol

- · Message sent to directory causes two actions:
 - Update the directory
 - More messages to satisfy request
- Block is in Uncached state: the copy in memory is the current value; only possible requests for that block are:
 - Read miss: requesting processor sent data from memory & requestor made <u>only</u> sharing node; state of block made Shared.
 - Write miss: requesting processor is sent the value & becomes the Sharing node. The block is made Exclusive to indicate that the only valid copy is cached. Sharers indicates the identity of the owner.
- Block is Shared ⇒ the memory value is up-to-date:
 - Read miss: requesting processor is sent back the data from memory & requesting processor is added to the sharing set.
 - Write miss: requesting processor is sent the value. All processors in the set Sharers are sent invalidate messages, & Sharers is set to identity of requesting processor. The state of the block is made Exclusive.

Example Directory Protocol

- Block is Exclusive: current value of the block is held in the cache of the processor identified by the set Sharers (the owner) ⇒ three possible directory requests:
 - Read miss: owner processor sent data fetch message, causing state of block in owner's cache to transition to Shared and causes owner to send data to directory, where it is written to memory & sent back to requesting processor.
 Identity of requesting processor is added to set Sharers, which still
 - contains the identity of the processor that was the owner (since it still has a readable copy). State is shared.

 Data write-back: owner processor is replacing the block and hence must write it back, making memory copy up-to-date (the home directory essentially becomes the owner), the block is now Uncached, and the Sharer set is empty.
 - Write miss: block has a new owner. A message is sent to old owner
 causing the cache to send the value of the block to the directory from
 which it is sent to the requesting processor, which becomes the new
 owner. Sharers is set to identity of new owner, and state of block is
 made Exclusive.

Example

Processor 1 Processor 2 Interconnect Directory Memory P1 P2 Bus Directory Memory Statel, Addt Valut, Statel, Addt Valut, Action Proc. Addt V

| | P1 | | | P2 | | | Bus | | | | Direc | ctory | | Memor |
|--------------------|-------|------|-------|-------|------|-------|--------|-------|------|-------|-------|-------|--------|-------|
| step | State | Addr | Value | State | Addi | Value | Action | Proc. | Addr | Value | Addi | State | {Procs | Value |
| P1: Write 10 to A1 | | | | | | | | | | | | | | |
| | | | | | | | | | | | | | | |
| P1: Read A1 | | | | | | | | | | | | | | |
| P2: Read A1 | | | | | | | | | | | | | | |
| | | | | | | | | | | | | | | |
| | | | | | | | | | | | | | | |
| P2: Write 20 to A1 | | | | | | | | | | | | | | |
| | | | | | | | | | | | | | | |
| P2: Write 40 to A2 | | | | | | | | | | | | | | |
| | | | | | | | | | | | | | | |
| | | | | | | | | | | | | | | |

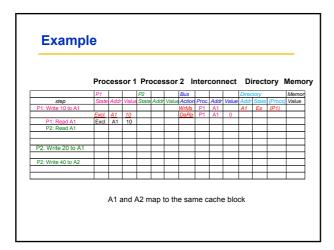
A1 and A2 map to the same cache block

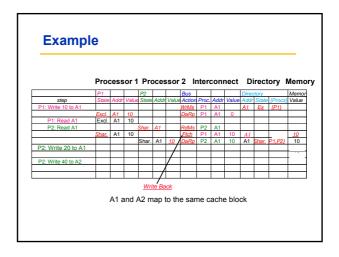
Example

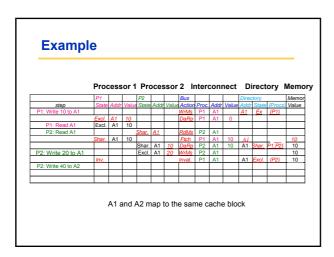
Processor 1 Processor 2 Interconnect Directory Memory

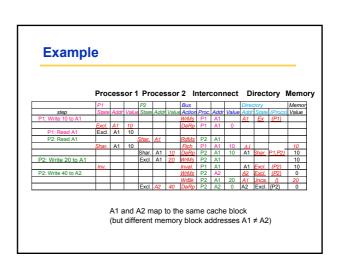
| | P1 | | | P2 | | | Bus | | | | Direc | | | Memor |
|--------------------|-------|------|-------|-------|------|-------|--------|-------|------|-------|-----------|-------|---------|-------|
| step | State | Addr | Value | State | Addr | Value | Action | Proc. | Addr | Value | Addr | State | (Procs) | Value |
| P1: Write 10 to A1 | | | | | | | WrMs | P1 | A1 | | <u>A1</u> | Ex | {P1} | |
| | Excl. | A1 | 10 | | | | DaRp | P1 | A1 | 0 | | | | |
| P1: Read A1 | | | | | | | | | | | | | | |
| P2: Read A1 | | | | | | | | | | | | | | |
| | | | | | | | | | | | | | | |
| | | | | | | | | | | | | | | |
| P2: Write 20 to A1 | | | | | | | | | | | | | | |
| | | | | | | | | | | | | | | |
| P2: Write 40 to A2 | | | | | | | | | | | | | | |
| | | | | | | | | | | | | | | |
| | | | | | | | | | | | | | | |
| | | | | | | | | | | | | | | |

A1 and A2 map to the same cache block



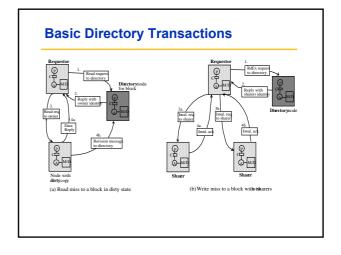


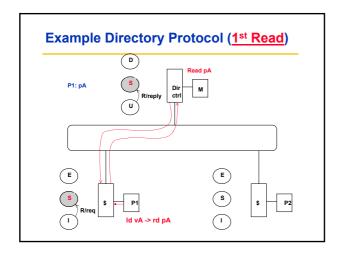


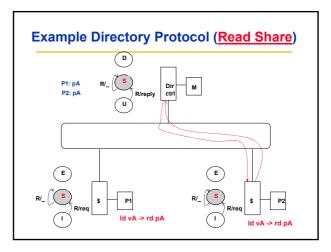


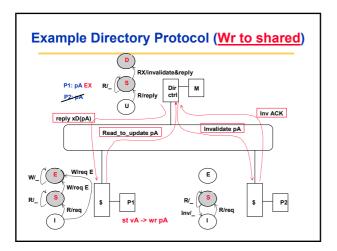
Implementing a Directory

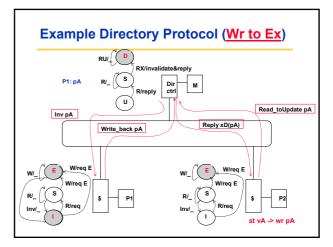
- We assume operations atomic, but they are not; reality is much harder; must avoid deadlock when run out of buffers in network (see Appendix E)
- · Optimizations:
 - read miss or write miss is Exclusive: send data directly to requestor from owner vs. first to memory and then from memory to requestor











A Popular Middle Ground

- · Two-level "hierarchy"
- Individual nodes are multiprocessors, connected non-hierarchically
 - e.g. mesh of SMPs
- Coherence across nodes is directory-based
 directory keeps track of nodes, not individual processors
- Coherence within nodes is snooping or directory
 orthogonal, but needs a good interface of functionality
- · SMP on a chip directory + snoop?

Synchronization

- Why Synchronize? Need to know when it is safe for different processes to use shared data
- Issues for Synchronization:
 - Uninterruptable instruction to fetch and update memory (atomic operation);
 - User level synchronization operation using this primitive;
 - For large scale MPs, synchronization can be a bottleneck; techniques to reduce contention and latency of synchronization

Uninterruptable Instruction to Fetch and Update Memory

Atomic exchange: interchange a value in a register for a value in memory

0 ⇒ synchronization variable is free

- $1 \Rightarrow$ synchronization variable is locked and unavailable
- Set register to 1 & swap
- New value in register determines success in getting lock
 0 if you succeeded in setting the lock (you were first)
 1 if other processor had already claimed access
- Key is that exchange operation is indivisible
- Test-and-set: tests a value and sets it if the value passes the test
- Fetch-and-increment: it returns the value of a memory location and atomically increments it
 - 0 ⇒ synchronization variable is free

Uninterruptable Instruction to Fetch and Update Memory

- · Hard to have read & write in 1 instruction: use 2 instead
- Load linked (or load locked) + store conditional
 - Load linked returns the initial value
 - Store conditional returns 1 if it succeeds (no other store to same memory location since preceding load) and 0 otherwise
- Example doing atomic swap with LL & SC:

R3,R4 R2,0(R1) R3,0(R1) ; mov exchange value ; load linked ; store conditional mov II ; branch store fails (R3 = 0) ; put load value in R4 beaz R3,try R4,R2

· Example doing fetch & increment with LL & SC:

; load linked ; increment (OK if reg-reg) ; store conditional ; branch store fails (R2 = 0) R2,0(R1) R2,R2,#1 addi beqz

User Level Synchronization— Operation Using this Primitive

Spin locks: processor continuously tries to acquire, spinning around a loop trying to get the lock

li R2,#1 exch R2,0(R1) bnez R2,lockit

- · What about MP with cache coherency?
 - Want to spin on cache copy to avoid full memory latency
 - Likely to get cache hits for such variables
- Problem: exchange includes a write, which invalidates all other copies; this generates considerable bus traffic
- Solution: start by simply repeatedly reading the variable; when it changes, then try exchange ("test and test&set"):

R2,#1 R3,0(R1) R3,lockit R2,0(R1) R2,try ;load var ;# 0 ⇒ not free ⇒ spin ;atomic exchange ;already locked?

Another MP Issue: Memory Consistency Models

What is consistency? When must a processor see the new value? e.g., seems that

A = 1; if (B == 0) ... B = 1; if (A == 0) .. L2:

- Impossible for both if statements L1 & L2 to be true?
 - What if write invalidate is delayed & processor continues?
- Memory consistency models: what are the rules for such cases?
- Sequential consistency: result of any execution is the same as if the accesses of each processor were kept in order and the accesses among different processors were interleaved ⇒ assignments before ifs above
 - SC: delay all memory accesses until all invalidates don

Memory Consistency Model

- · Schemes faster execution to sequential consistency
- Not an issue for most programs; they are synchronized
 - A program is synchronized if all access to shared data are ordered by synchronization operations write (x)

release (s) {unlock} acquire (s) {lock}

read(x)

- Only those programs willing to be nondeterministic are not synchronized: "data race": outcome f(proc. speed)
- Several Relaxed Models for Memory Consistency since most programs are synchronized; characterized by their attitude towards: RAR, WAR, RAW, WAW to different

Relaxed Consistency Models: The Basics

- Key idea: allow reads and writes to complete out of order, but to use synchronization operations to enforce ordering, so that a synchronized program behaves as if the processor were sequentially consistent

- By relaxing orderings, may obtain performance advantages
 Also specifies range of legal compiler optimizations on shared data
 Unless synchronization points are clearly defined and programs are
 synchronized, compiler could not interchange read and write of 2 shared
 data items because might affect the semantics of the program.
- 3 major sets of relaxed orderings:

 1. W—R ordering (all writes completed before next read)

 Because retains ordering among writes, many programs that operate under sequential consistency operate under this model, without additional synchronization. Called processor consistency.
- 2. W → W ordering (all writes completed before next write)
 3. R → W and R → R orderings, a variety of models depending on ordering restrictions and how synchronization operations enforce ordering
- Many complexities in relaxed consistency models; defining precisely what it means for a write to complete; deciding when processors can see values that it has written

Mark Hill observation

Instead, use speculation to hide latency from strict consistency model

- If processor receives invalidation for memory reference before it is committed, processor uses speculation recovery to back out computation and restart with invalidated memory reference
- Aggressive implementation of sequential consistency or processor consistency gains most of advantage of more relaxed models
- 2. Implementation adds little to implementation cost of speculative processor
- 3. Allows the programmer to reason using the simpler programming models

Cross Cutting Issues: Performance Measurement of Parallel Processors

- · Performance: how well scale as increase Proc
- · Speedup fixed as well as scaleup of problem
 - Assume benchmark of size n on p processors makes sense: how scale benchmark to run on m * p processors?
 - Memory-constrained scaling: keeping the amount of memory used per processor constant
 - <u>Time-constrained scaling</u>: keeping total execution time, assuming perfect speedup, constant
- Example: 1 hour on 10 P, time ~ O(n3), 100 P?
 - <u>Time-constrained scaling</u>: 1 hour ⇒ 10^{1/3}n ⇒ 2.15n scale up
 - Memory-constrained scaling: 10n size ⇒ 10³/10 ⇒ 100X or 100 hours! 10X processors for 100X longer???
 - Need to know application well to scale: # iterations, error tolerance

Fallacy: Amdahl's Law doesn't apply to parallel computers

- · Since some part linear, can't go 100X?
- · 1987 claim to break it, since 1000X speedup
 - researchers scaled the benchmark to have a data set size that is 1000 times larger and compared the uniprocessor and parallel execution times of the scaled benchmark. For this particular algorithm the sequential portion of the program was constant independent of the size of the input, and the rest was fully parallel—hence, linear speedup with 1000 processors
- · Usually sequential scale with data too

Fallacy: Linear speedups are needed to make multiprocessors cost-effective

Mark Hill & David Wood 1995 study

- · Compare costs SGI uniprocessor and MP
 - Uniprocessor = \$38,400 + \$100 * MB
 - MP = \$81,600 + \$20,000 * P + \$100 * MB
- 1 GB: uni = \$138k v. mp = \$181k + \$20k * P
- What speedup for better MP cost performance?
 - 8 proc = \$341k; \$341k/138k ⇒ 2.5X
 - 16 proc ⇒ need only 3.6X, or 25% linear speedup
- Even if need some more memory for MP, not linear

Fallacy: Scalability is almost free

- "build scalability into a multiprocessor and then simply offer the multiprocessor at any point on the scale from a small number of processors to a large number"
- Cray T3E scales to 2048 CPUs vs. 4 CPU Alpha
 - At 128 CPUs, it delivers a peak bisection BW of 38.4 GB/s, or 300 MB/s per CPU (uses Alpha microprocessor)
 - Compaq Alphaserver ES40 up to 4 CPUs and has 5.6 GB/s of interconnect BW, or 1400 MB/s per CPU
- Build apps that scale requires significantly more attention to load balance, locality, potential contention, and serial (or partly parallel) portions of program. 10X is very hard

Pitfall: Not developing SW to take advantage (or optimize for) multiprocessor architecture

- SGI OS protects the page table data structure with a single lock, assuming that page allocation is infrequent
- Suppose a program uses a large number of pages that are initialized at start-up
- Program parallelized so that multiple processes allocate the pages
- But page allocation requires lock of page table data structure, so even an OS kernel that allows multiple threads will be serialized at initialization (even if separate processes)

Answers to 1995 Questions about Parallelism

- In the 1995 edition of this text, we concluded the chapter with a discussion of two then current controversial issues.
- What architecture would very large scale, microprocessor-based multiprocessors use?
- 2. What was the role for multiprocessing in the future of microprocessor architecture
- Answer 1. Large scale multiprocessors did not become a major and growing market ⇒ clusters of single microprocessors or moderate SMPs
- Answer 2. Astonishingly clear. For at least for the next 5 years, future MPU performance comes from the exploitation of TLP through multicore processors vs. exploiting more ILP

Cautionary Tale

- Key to success of birth and development of ILP in 1980s and 1990s was software in the form of optimizing compilers that could exploit ILP
- Similarly, successful exploitation of TLP will depend as much on the development of suitable software systems as it will on the contributions of computer architects
- · Given the slow progress on parallel software in the past 30+ years, it is likely that exploiting TLP broadly will remain challenging for years to come

T1 ("Niagara")

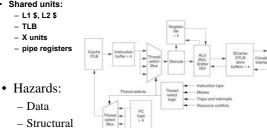
- **Target: Commercial server applications**
 - High thread level parallelism (TLP)
 - » Large numbers of parallel client requests
 - Low instruction level parallelism (ILP)
 - » High cache miss rates
 - » Many unpredictable branches
 - » Frequent load-load dependencies
- Power, cooling, and space are major concerns for data centers
- Metric: Performance/Watt/Sq. Ft.
- Approach: Multicore, Fine-grain multithreading, Simple pipeline, Small L1 caches, Shared L2



T1 Architecture · Also ships with 6 or 4 processors ating Point Unit DRAM Control Sparc 0 DDR2 L2 Channel 0 Sparc 1 DDR2 Sparc 2 Sparc 3 DDR2 nel 2 Sparc 5 Channel 3 Sparc 6 L2 Bank 3 Sparc 7 Contro JTAG Clock & Test JBUS System JBUS (200 MHz) Unit SSI (50 MHz)

T1 pipeline

- Single issue, in-order, 6-deep pipeline: F, S, D, E, M, W
- 3 clock delays for loads & branches.
- Shared units:

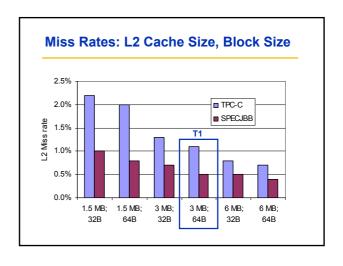


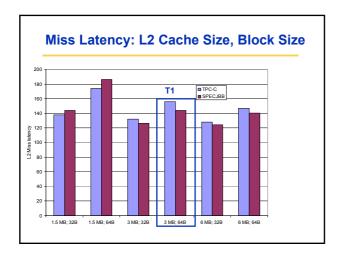
T1 Fine-Grained Multithreading

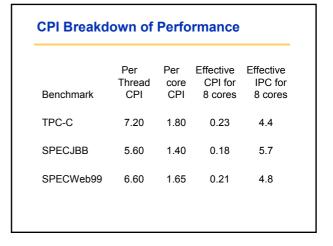
- Each core supports four threads and has its own level one caches (16KB for instructions and 8 KB for data)
- · Switching to a new thread on each clock cycle
- · Idle threads are bypassed in the scheduling
 - Waiting due to a pipeline delay or cache miss
 - Processor is idle only when all 4 threads are idle or stalled
- · Both loads and branches incur a 3 cycle delay that can only be hidden by other threads
- A single set of floating point functional units is shared by all 8 cores
 - floating point performance was not a focus for T1

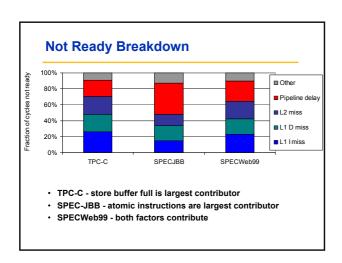
Memory, Clock, Power 16 KB 4 way set assoc. I\$/ core 18 KB 4 way set assoc. D\$/ core 18 KB 4 way set assoc. L2 \$ shared 19 4 x 750KB independent banks 19 crossbar switch to connect 10 crycle throughput, 8 cycle latency 10 Direct link to DRAM & Jbus 10 Manages cache coherence for the 8 cores 10 CAM based directory 10 Coherency is enforced among the L1 caches by a directory associated with each L2 cache block 10 Used to track which L1 caches have copies of an L2 block 10 By associating each L2 with a particular memory bank and enforcing the subset property. T1 can place the directory at L2 rather than at the memory, which reduces the directory overhead 11 data cache is write-through, only invalidation messages are required; the data can always be retrieved from the L2 cache

1.2 GHz at ≈72W typical, 79W peak power consumption









| Performance: E | Benchr | narks + | Sun | Marketing | |
|---|-------------------|-----------------|-----|--|--|
| Benchmark\Architecture | Sun Fire T2000 | IBM p5-550 v | | Dell PowerEdge | |
| SPECjbb2005 (Java server software) business operations/ sec | 63,378 | 61,789 | | 24,208 (SC1425 with dual single core Xeon) | |
| SPECweb2005 (Web server performance) | 14,001 | 7,881 14,740 | | 4,850 (2850 with two dual-core Xeon processors) | |
| NotesBench (Lotus Notes performance) | 16,061 | | | | |
| SPECjappServer 2004 Dual Node | | | | | |
| | Sun Fire T2000 | HP rx4640 | | | |
| Space (RU) | 2 | - 4 | | | |
| Watts | 320 | 1,303 | | | |
| Performance (SPECjapp JOPs) | 615 | 471 | | | |
| Performance / Watt | 1.922 | 0.361 | C | ce, Watts, and Performance | |
| | | | | | |

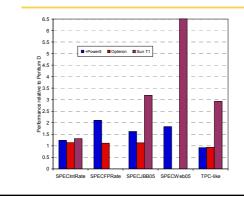
HP marketing view of T1 Niagara

- Sun's radical UltraSPARC T1 chip is made up of individual cores that have much slower single thread performance when compared to the higher performing cores of the Intel Xeon, Itanium, AMD Opteron or even classic UltraSPARC
- The Sun Fire T2000 has poor floating-point performance, by Sun's own admission.
- The Sun Fire T2000 does not support commerical Linux or Windows® and requires a lock-in to Sun and Solaris.
- The UltraSPARC T1, aka CoolThreads, is new and unproven, having just been introduced in December 2005.
- In January 2006, a well-known financial analyst downgraded Sun on concerns over the UltraSPARC T1's limitation to only the Solaris operating system, unique requirements, and longer adoption cycle, among other things. [10]
- Where is the compelling value to warrant taking such a risk?
- http://h71028.www7.hp.com/ERC/cache/280124-0-0-0-121.html

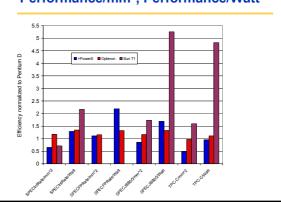
| Microprocessor | Comparison |
|----------------|------------|
|----------------|------------|

| Processor | SUN T1 | Opteron | Pentium D | IBM Power 5 |
|-----------------------------------|------------------|---------------|----------------|------------------|
| Cores | 8 | 2 | 2 | 2 |
| Instruction issues / clock / core | 1 | 3 | 3 | 4 |
| Peak instr. issues | 8 | 6 | 6 | 8 |
| Multithreading | Fine- grained | No | SMT | SMT |
| L1 I/D in KB per core | 16/8 | 64/64 | 12K uops/16 | 64/32 |
| L2 per core/shared | 3 MB shared | 1MB / core | 1MB/ core | 1.9 MB shared |
| Clock rate (GHz) | 1.2 | 2.4 | 3.2 | 1.9 |
| Transistor count (M) | 300 | 233 | 230 | 276 |
| Die size (mm²) | 379 | 199 | 206 | 389 |
| Power (W) | 79 | 110 | 130 | 125 |

Performance Relative to Pentium D



Performance/mm², Performance/Watt



Niagara 2

- Improve performance by increasing threads supported per chip from 32 to 64
 - 8 cores * 8 threads per core
- · Floating-point unit for each core, not for each
- Hardware support for encryption standards EAS, 3DES, and elliptical-curve cryptography
- Niagara 2 will add a number of 8x PCI Express interfaces directly into the chip in addition to integrated 10Gigabit Ethernet XAU interfaces and Gigabit Ethernet ports.
- Integrated memory controllers will shift support from DDR2 to FB-DIMMs and double the maximum amount of system memory.

"Sun's Niagara Begins CMT Flood
The Sun UltraSPARC T1 Processor Released'
Microprocessor Report, January 3, 2006

And in Conclusion ...

- Caches contain all information on state of cached memory blocks
- Snooping cache over shared medium for smaller MP by invalidating other cached copies on write
- Sharing cached data ⇒ Coherence (values returned by a read), Consistency (when a written value will be returned by a read)
- Snooping and Directory Protocols similar; bus makes snooping easier because of broadcast (snooping ⇒ uniform memory access)
- Directory has extra data structure to keep track of state of all cache blocks
- Distributing directory
 ⇒ scalable shared address multiprocessor
- ⇒ Cache coherent, Non uniform memory access

Reading

- This lecture:
 - chapter 4: 4.4-4.10 rest of Multiprocessors and TLP
- · Next lecture:
 - chapter 5: Advanced Memory Hierarchy