

# Advanced Cache Topics

# Cache Coherency & Consistency

- A must, to avoid using wrong data (e.g., MP, DMA)
- Solution:  
Pass information about changes between caches & memory
- The process in which one element notifies of its actions and others listen is called SNOOPING \*
- Idea:
  - ◆ Main memory is passive
  - ◆ When needed, caches distribute changes to memory or other caches
  - ◆ All caches listen to snoop messages and act when needed

\* Other schemes exist, but SNOOPing is the only practical scheme in use

# Write Through Simple Coherency Protocol

- All memory writes are written to main memory, even if the line is in cache
- Each line has 2 states: valid/invalid
  - ◆ A line is invalidated if another processor changed the memory portion occupied by that line
- Activity:

Action	Cache	Bus	Other Caches
Read Hit	Read	---	---
Read Miss	Read	Line Fill	---
Write Hit	Write	Write	Snoop/INV
Write Miss	---	Write	Snoop/INV

 Simple

 Slow, lot of bus traffic (immediate writes, small chunks)

- Write buffers are used to avoid write latency issues

# Write Back

- Distributed cache management
- Idea: Hold a line in the cache, even if modified, as long as other agents do not need it.
- Update content externally, only when another agent needs it
- “Agent” informs others of intentions:
  - ◆ Read a copy
  - ◆ Modify a copy
  - ◆ Broadcast modifications
- Others respond
  - ◆ Have a copy
  - ◆ Have a dirty copy
- Most common model: MESI (Modified, Exclusive, Shared, Invalid)

# MESI

## ● Assumptions

- ◆ Data ownership model. Only one cache can have dirty data
- ◆ Writes are typically not broadcasted (i.e. *data* is not transmitted). Cause data invalidation in other caches

## ● Bus activity

- ◆ Master: “Intend to cache” (inquiry cycle)
- ◆ Slave: “I have a copy” (Hit signal)  
“I have a modified copy” (HitM signal)

## ● Each cache line can be in one of the following states:

- ◆ Modified: line resides exclusively in this cache only  
Content is modified relative to memory
- ◆ Exclusive: line resides exclusively in this cache only  
Content is same as memory
- ◆ Shared: line resides in this cache but may be shared with other  
Content is same as memory
- ◆ Invalid: Line contains no valid memory copy

## ● The idea: modified/exclusive lines are “owned” by the cache They can be changed w/o telling other caches

## ● Attempt to access a non-owned line should be broadcasted

- ◆ If owned by another: the owner should update the world and give up ownership

# MESI Cache Line States

In the data cache, a cache protocol known as MESI maintains consistency with caches of other processors and with an external cache. The data cache has two status bits per tag; so, each line can be in one of the states defined in Table 18. The state of a cache line can change as the result of either internal or external activity related to that line. In general, the operation of the MESI protocol is transparent to programs.

<b>Cache Line State:</b>	<b>M Modified</b>	<b>E Exclusive</b>	<b>S Shared</b>	<b>I Invalid</b>
This cache line is valid?	Yes	Yes	Yes	No
The memory copy is...	...out of date	...valid	...valid	—
Copies exist in caches of other processors?	No	No	Maybe	Maybe
A write to this line ...	...does not go to bus	...does not go to bus	...goes to bus and updates cache	...goes directly to bus

**Table 18-1. MESI Cache Line States**

# MESI

Action	Current State	Next State	Bus Activity
<hr/>			
Read	M	M	None (Read Hit)
Read	E	E	None (Read Hit)
Read	S	S	None (Read Hit)
Read	I	S / E	Line Fill (Read Miss) - Send Inq (intend to read) to other - Bring data - Implementation may decide if S / E *
Write	M	M	None (Write Hit)
Write	E	M	None (Write Hit)
Write	S	S / E	Write Through (Write Hit) - Send INV signal to other caches
Write	I	I	Write Through (Write Miss) - Send INV signal - Wait for line to be written by other - If write allocate then “read for ownership” (to E stage)

\* e.g., S may be better if line exist in another cache, otherwise E

# MESI (Cont)

Action	Current State	Next State	Bus Activity
Snoop Read	M	S	Write Back - Send Hit & HitM
Snoop Read	E	S	- Send Hit
Snoop Read	S	S	- Send Hit
Snoop Read	I	I	Do nothing
Snoop Write	M	I	Write Back - Send Hit/HitM - Invalidate owned copy
Snoop Write	E	I	Invalidate owned copy
Snoop Write	S	I	Invalidate owned copy
Snoop Write	I	I	Invalidate owned copy



# MESI (Cont)

## Variations (*not* Pentium):

### ● Ownership Exchange

- ◆ Allows Modified line to be exchanged between caches without write-back

### ● Write Allocate

- ◆ Write miss causes line fill (“read for ownership”), followed as write hit to E line.
- + Higher performance if line is used later
- Reduced if line is not used (e.g. zeroing big array)
- Must identify non-cacheable memory (e.g., video RAM)

### ● Handling 2 level caches is more complex. See later...

# X86 Caches

Processor	L1	L2	Comments
386	---	32K	82385
486	8K I+D	?	
Pentium®	8K I + 8K D	256K/512K	82491/82496
Pentium® MMX	16K I + 16K D	256K/512K	82491/82496
Pentium Pro®	8K I + 8K D	256K-1024K	On package
Pentium-II®	16K I + 16K D	512K	On SECC

## ● Pentium® internal caches

- ◆ Data: 8K, 128 sets, 32b line, 2 Way, LRU, Write back/MESI (2 bits)
- ◆ Inst: 8K, 128 sets, 32b line, 2 Way, LRU, I/V bits (1 bit)

## ● 3 port tags, 1 port data, banked cache

- ◆ Data: 1 snoop, 2 U/V pipes
- ◆ Inst: 1 snoop, 2 prefetch

# Processor-Initiated Read Cycles

Present State	Pin Activity	Next State	Description
M	n/a	M	Read hit; data is provided to processor core by cache. No bus cycle is generated.
E	n/a	E	Read hit; data is provided to processor core by cache. No bus cycle is generated.
S	n/a	S	Read hit; data is provided to the processor by the cache. No bus cycle is generated.
I	CACHE# low AND KEN# low AND WB/WT# high AND PWT low	E	Data item does not exist in cache (MISS). A bus cycle (read) will be generated by the Pentium® processor. This state transition will happen if WB/WT# is sampled high with first BRDY# or NA#.
I	CACHE# low AND KEN# low AND (WB/WT# low OR PWT high)	S	Same as previous read miss case except that WB/WT# is sampled low with first BRDY# or NA#.
I	CACHE# high OR KEN# high	I	KEN# pin inactive; the line is not intended to be cached in the Pentium processor.

**NOTE:** \*Locked accesses to the data cache will cause the accessed line to transition to the Invalid state

The transition from I to E or S states (based on WB/WT#) happens only if KEN# is sampled low with the first of BRDY# or NA#, and the cycle is transformed into a LINE FILL cycle. If KEN# is sampled high, the line is not cached and remains in the I state.

**Table 2-4. Data Cache State Transitions for UNLOCKED Pentium® Processor Initiated Read Cycles\***

# Processor-Initiated Write Cycles

Present State	Pin Activity	Next State	Description
M	n/a	M	Write hit; update data cache. No bus cycle generated to update memory.
E	n/a	M	Write hit; update cache only. No bus cycle generated; line is now MODIFIED.
S	PWT low AND WB/WT# high	E	Write hit; data cache updated with write data item. A write-through cycle is generated on bus to update memory and/or invalidate contents of other caches. The state transition occurs after the writethrough cycle completes on the bus (with the last BRDY#).
S	PWT low AND WB/WT# low	S	Same as above case of write to S-state line except that WB/WT# is sampled low.
S	PWT high	S	Same as above cases of writes to S state lines except that this is a write hit to a line in a writethrough page; status of WB/WT# pin is ignored.
I	n/a	I	Write MISS; a writethrough cycle is generated on the bus to update external memory. No allocation done.

**NOTE:** Memory writes are buffered while I/O writes are not. There is no guarantee of synchronization between completion of memory writes on the bus and instruction execution after the write. A serializing instruction needs to be executed to synchronize writes with the next instruction if necessary.

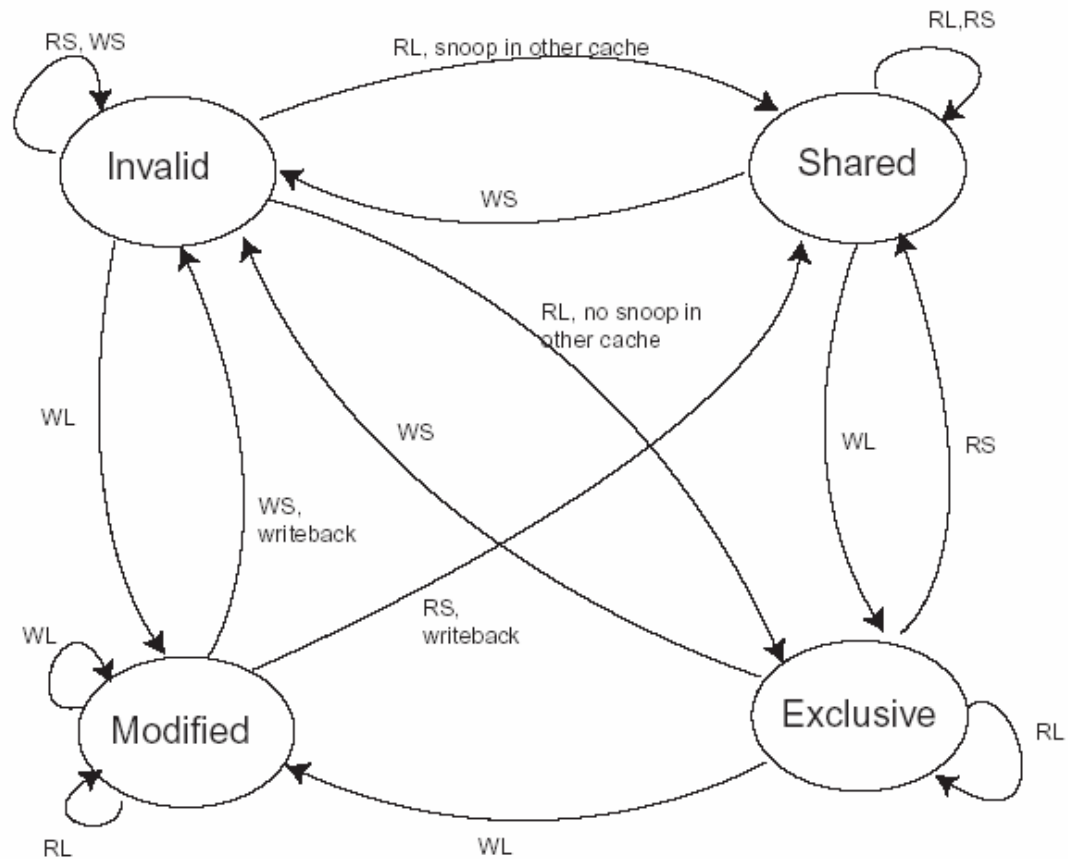
**Table 2-5. Data Cache State Transitions for Pentium® Processor Initiated Write Cycles**

# Cache State Transitions During Inquire Cycles

Present State	Next State INV=1	Next State INV=0	Description
M	I	S	Snoop hit to a MODIFIED line indicated by HIT# and HITM# pins low. Pentium ® processor schedules the writing back of the modified line to memory.
E	I	S	Snoop hit indicated by HIT# pin low; no bus cycle generated.
S	I	S	Snoop hit indicated by HIT# pin low; no bus cycle generated.
I	I	I	Address not in cache; HIT# pin high.

**Table 2-6. Cache State Transitions During Inquire Cycles**

# MESI Diagram

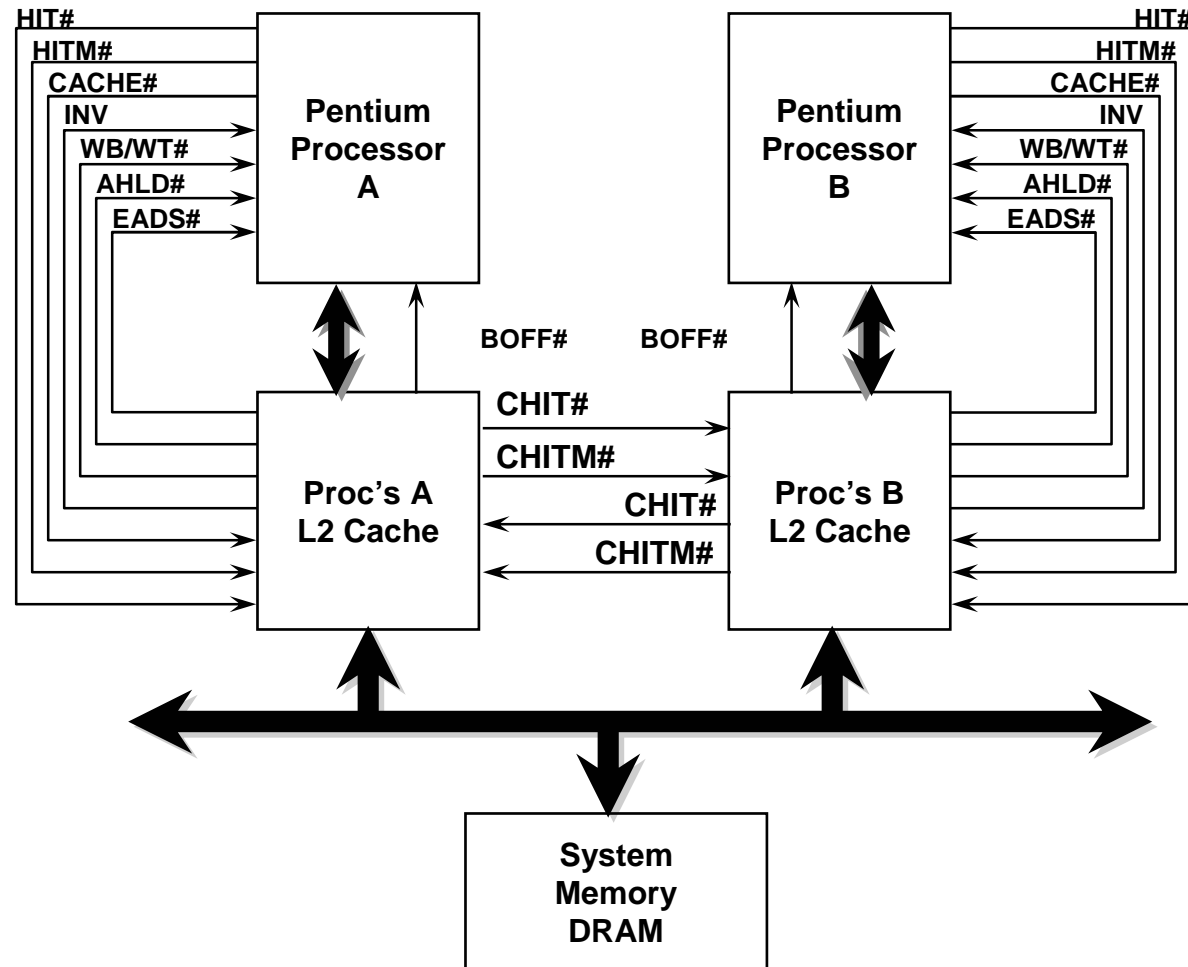


Write Snoop = WS  
Read Snoop = RS  
Write Local CPU = WL  
Read Local CPU = RL

# Cache Relevant Signals

<b>AHOLD</b>	<b>I</b>	<b>Address Hold</b>
<b>BOFF#</b>	<b>I</b>	<b>Start the aborted bus cycle from the beginning</b>
<b>CACHE#</b>	<b>O</b>	<b>Internal cacheability of the cycle</b>
<b>EADS#</b>	<b>I</b>	<b>Valid external address driven on address pins</b>
<b>EWBE#</b>	<b>I</b>	<b>External Write Buffers Empty (sync info)</b>
<b>FLUSH#</b>	<b>I</b>	<b>Force writeback of all modified lines</b>
<b>HIT#</b>	<b>O</b>	<b>Hit answer following an Inquire request</b>
<b>HITM#</b>	<b>O</b>	<b>Hit modify answer following an Inquire req.</b>
<b>INV</b>	<b>I</b>	<b>Inquire/Invalidate request</b>
<b>KEN#</b>	<b>I</b>	<b>Cache enable pin</b>
<b>PCD/PWT</b>	<b>O</b>	<b>Page Cache Disable / Page Write Through</b>
<b>WB/WT#</b>	<b>I</b>	<b>Write Back/Write Through cache line request</b>

# Pentium System - 2 Levels of Caches



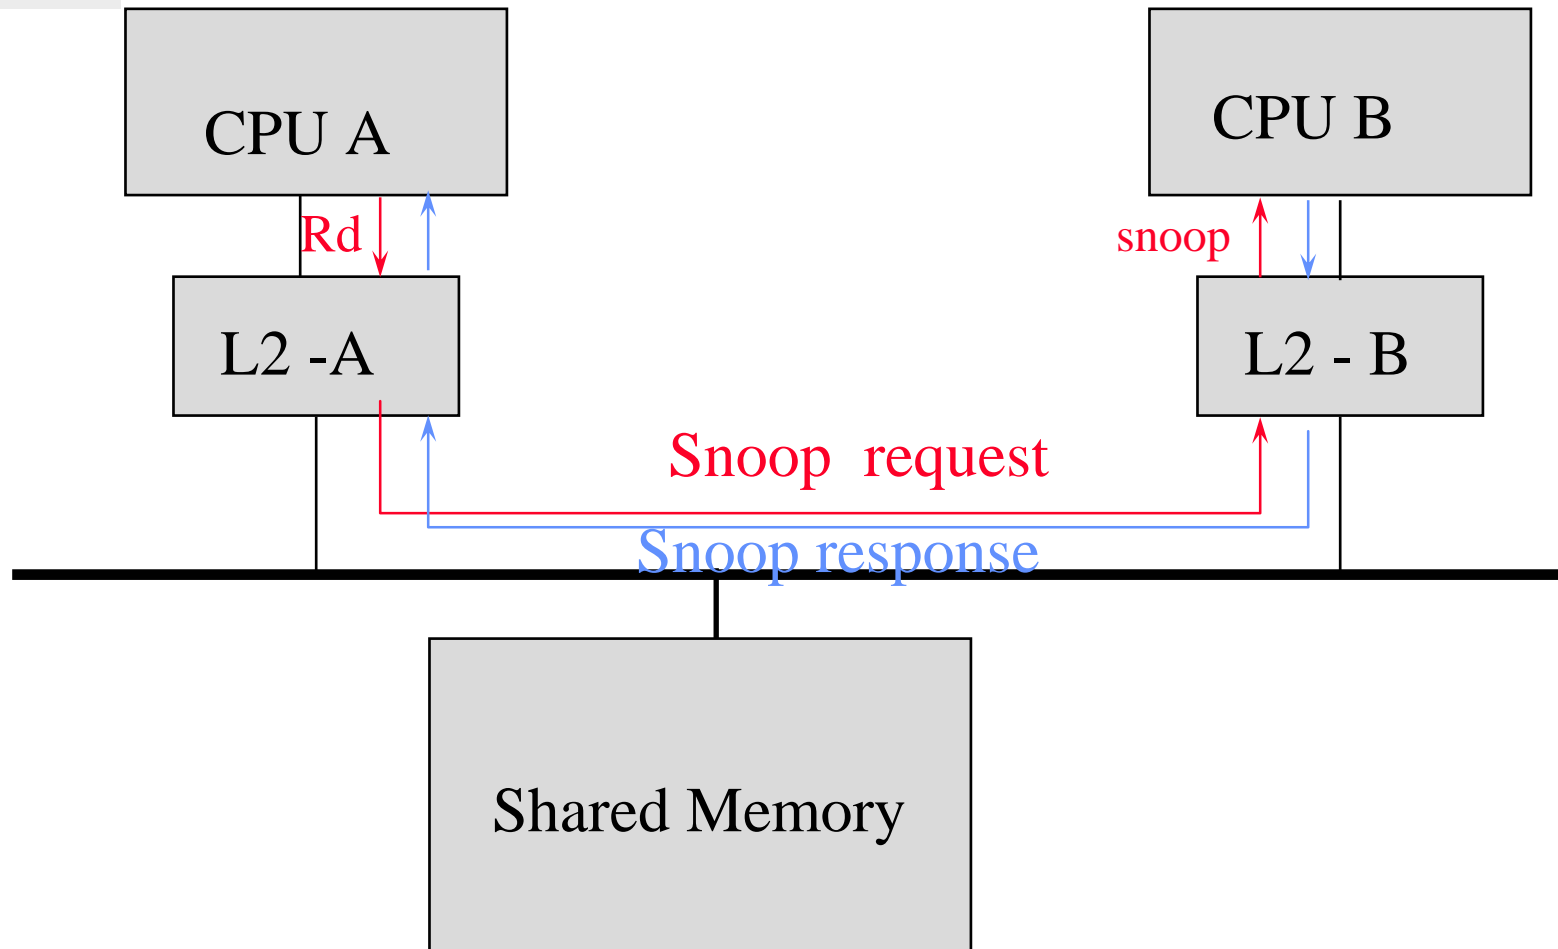
Two Processors With Writeback Caches Using the MESI Model





# L1/L2 organization

## “Flow” of SNOOP cycles



# L1/L2 communication

Case	L2	L1
<hr/>		
1 First read, No partners	E	S
2 First write to line	M	E
3 Subsequent writes	M	M
4 Another processor read to same line	S	S

- Each line in L1 must be in L2 - the *inclusion property*
- L2 status represents the combined L1+L2 complex
- L1 status is relative to its L2 and always  $L1 \leq L2$  (I,S,E,M=0,1,2,3)
  - ◆ Every CPU *M* or *E* line (dirty or exclusive) is *M* in L2.
  - ◆ Every L2 *M* line might be *I*, *S*, *E*, *M* in L1.
- All snoops into L2 are passed to L1
  - ◆ L2 does not really know L1 status!
  - ◆ But L2 should wait for L1 only if L2 is in M state

# Process Synchronization

- Distinguish between *memory coherency* to *synchronization*
  - ◆ “If 2 processors increment the same cell, how to ensure  $X \leftarrow X+2$ ”
  - ◆ This is a synchronization problem, not a cache issue!
- Caches do not guarantee order!  
**Semaphores** should be used for exclusive access to X!
- In X86, XCHG can be used to implement a semaphore
  - ◆ XCHG has an implicit lock which makes it an atomic operation
  - ◆ Alternative: use explicit LOCK prefix on the instruction

```
get_x:
    MOV     EAX, 0
again:  XCHG  EAX, x_guard    / atomic - locked
    CMP     EAX, 0
    JZ      again          / until success
    .
    .
    MOV     EAX, X          / manipulate value
    INC     EAX             / ensure X is untouched
    MOV     X, EAX          / between operations.
    .
free_x: MOV     x_guard, 1
```